

Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

Davide Fissore  

Université Côte d’Azur, Inria, France

Enrico Tassi  

Université Côte d’Azur, Inria, France

Abstract

We mechanize two operational semantics for PROLOG with cut and prove them equivalent, in the Rocq prover. The first semantics closely models the ELPI’s implementation, it is based on a stack of choice points and is well suited for studying optimizations employed by PROLOG abstract machines. The second semantics is instead based on and-or trees, in which the branches pruned by the cut operator are made explicit.

We use the tree-based semantics (1) to prove properties about the effect of the cut operator in the evaluation of choice points, since the rules from classical logic do not apply, and (2) to reason about the determinacy analysis introduced in Elpi version 3.0, which statically identifies computations that produce at most one result.

2012 ACM Subject Classification Theory of computation → Operational semantics

Keywords and phrases Semantics, Prolog, Elpi, Determinacy Analysis, Rocq, Coq

Digital Object Identifier 10.4230/LIPIcs.CVIT.2016.23

Funding *Davide Fissore*: This work has been supported by the French government, through the France 2030 investment plan managed by the Agence Nationale de la Recherche, as part of the “UCA DS4H” project, reference ANR-17-EURE-0004.

1 Introduction

ELPI is a dialect of λ PROLOG (see [14, 15, 7, 12]) used as an extension language for the ROCQ prover (formerly the COQ proof assistant). ELPI has become an important infrastructure component: several projects and libraries depend on it [13, 3, 4, 19, 8, 9]. Examples include the Hierarchy-Builder library-structuring tool [5] and Derive [17, 18, 11], a program-and-proof synthesis framework with industrial applications at SkyLabs AI.

Starting with version 3, ELPI is equipped with a static analysis for determinacy [10] to help users tame backtracking. ROCQ users are familiar with functional programming but not necessarily with logic programming, in particular uncontrolled backtracking is a common source of inefficiency and makes debugging harder. The determinacy checker identifies predicates that behave like functions, i.e., predicates that commit to their first solution and leave no *choice points* (places where backtracking could resume).

This paper reports our first steps towards a mechanization, in the ROCQ prover, of the determinacy analysis from [10]. We focus on the control operator *cut*, which is useful to restrict backtracking but makes the semantics depart from a pure logical reading.

We formalize two operational semantics for PROLOG with cut. The first is a stack-based semantics that closely models ELPI’s implementation and is similar to the semantics mechanized by Pusch in ISABELLE/HOL [16] and to the model of Debray and Mishra [6, Sec. 4.3]. This stack-based semantics is a good starting point to study further optimizations used by standard PROLOG abstract machines [20, 1], but it makes reasoning about the scope of cut difficult. To address this limitation we introduce a tree-based semantics in which the branches pruned by *cut* are explicit and we prove the two semantics equivalent. Using the



© Jane Open Access and Joan R. Public;
licensed under Creative Commons License CC-BY 4.0

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:14

Leibniz International Proceedings in Informatics



LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

```

Inductive P := IP of nat. Inductive D := ID of nat. Inductive V := IV of nat.

Inductive Tm :=
  | Tm_P of P
  | Tm_D of D
  | Tm_V of V
  | Tm_App of Tm & Tm.

Record Unif := {
  unify : Tm → Tm → Σ → option Σ;
  matching : Tm → Tm → Σ → option Σ;
}.

```

■ Figure 1 Tm types

■ Figure 2 Unif type

tree-based semantics, we then prove some properties about the impact of evaluating the *cut* in query disjunctions. Unlike in classical logic, $q_1 \vee q_2$ is not equivalent to $q_2 \vee q_1$, since q_2 may be cut-away by q_1 . Moreover, we show that if every rule of a predicate passes the determinacy analysis, then a call to that predicate does not leave any choice points.

2 Preliminaries: syntax and informal semantics

Before going to the two semantics, we show the piece of data structure that are shared by the them. The smallest unit of code that we can use in the language is an atom. The atom inductive (see Type 1) is either a cut or a call. A call carries a term (see Figure 1). A term (Tm) is either a predicate, a datum, a variable or the binary application of two terms.

```

Inductive A := cut | call : Tm → A. (1)
Record R := mkR { head : Tm; premises : list A }. (2)
Record P := { rules : seq R; sig : sigT }. (3)
Definition Σ := {fmap V → Tm}. (4)
Definition bc : P →  $\mathcal{F}_V$  → Tm → Σ →  $\mathcal{F}_V$  * seq (Σ * seq A) := (5)

```

A rule (see Type 2) has a head of type term and a list of premises, that are a list of atoms. A program (see Type 3) is made by its rules and the signatures of predicates, **sigT** is a mapping from predicates to their signatures, signatures represent the definition of type from the simply typed lambda calculus, i.e. it is either a base type or an arrow for term application. We decorate arrows to know the mode (input or output) of the application argument.

A substitution (see Type 4) is a mapping from variables to terms. It is the output of a successful query. The empty substitution is noted ε .

The backchain function (bc, see Type 5) takes: a program P ; a set fv of variables ($\mathcal{F}_V$); a query q ; and the substitution σ . A backchain is performed when the current atom is a **call**. In this case, it start by recursively *dereferencing* the variables in σ , i.e. replacing all the variables in the query by their term mapping in σ .

TODO: free variables and program freshness

Then it takes the head of the term, which should be rigid, i.e. a predicate p , and look for the signature s of p in the program. Then, for each rule r in P , it unifies the head of r with q . Unification is possible thanks to the unificator U , whose type is shown in Figure 2. U explains how to unify terms up to standard unification (for output terms) or matching (for input terms). The function *unif* (resp. *matching*) unifies (resp. matches) the two terms under σ , it return an extension of σ if the two terms can be unified (resp. matched), and None otherwise. The result of a backchain is made by filtering the rules in P that can be used on the given query. In particular, for each rule filtered rules r , it returns a pair containing the extension of σ making the head of r and q unify and the list of premises of r .

```

f 1 2.    f 2 3.
r 0 0.    r 0 1.    r 0 2 :- !.
g X Z :- r X Y, f Y Z, !.    % r1
g X Z :- f X Y, f Y Z.      % r2

```

■ **Figure 3** Small ELPI program example

```

Inductive tree :=
| KO | OK | TA of A
| Or of option tree & Σ & tree
| And of tree & seq A & tree.

```

■ **Figure 4** Tree inductive

2.1 The cut operator

ELPI is a dialect of λ PROLOG with the hard cut operator and, in the language we have shown before, *cut* is an atom. The semantics of the cut operator adopted in the ELPI language corresponds to the *hard cut* operator of standard SWI-PROLOG. This operator has two primary purposes. First, it eliminates all alternatives that are created either simultaneously with, or after, the introduction of the cut into the execution state.

The ELPI program we will use as an example in the following sections is shown in Figure 3.

To illustrate the high-level description of *cut*, consider the program P shown in Figure 3 and the query $q = g \ 0 \ R$. Both rules for g can be used on the query q . Backchaining q in P from the empty (ε) substitution returns the list $[(\sigma_1, [r \ X \ Y, f \ Y \ Z, !]), (\sigma_2, [f \ X \ Y, f \ Y \ Z])]$ with σ_1 and σ_2 both equal to $\{X \mapsto 0, Z \mapsto R\}$. The two elements of the list are generated respectively from the rules r_1 and r_2 .

The next step of the interpretation corresponds to the backchaining of $r \ X \ Y$ from σ_1 . This returns $[(\sigma_3, []), (\sigma_4, []), (\sigma_5, [])]$ with $\sigma_3 := \sigma_1 + \{Y \mapsto 0\}$, $\sigma_4 := \sigma_1 + \{Y \mapsto 1\}$ and $\sigma_5 := \sigma_1 + \{Y \mapsto 2\}$. The three lists of premises are empty since all rules for r have no premises.

Backchaining $f \ Y \ Z$ in σ_3 gives no solution, therefore, by backtracking, we backchain $f \ Y \ Z$ in σ_4 . This gives a solution, namely $\sigma_6 := \sigma_4 + \{R \mapsto 2\}$.

We now reach the cut operator and can see its double side behavior. In the current execution state, we have still two unexplored alternatives: we can backtrack on the third solution for r that assign Y to 2 and on the rules r_2 . We can note, however, the first alternative have been generated after the cut introduction, and the second alternative have been generated at the same moment of the cut. Therefore both are cut away leaving no choice point. The final complete solution is: $\{X \mapsto 0, Z \mapsto R, Y \mapsto 1, R \mapsto 2\}$, where X , Y and Z can be ignored since they are the result internal unification.

We propose two formal, operational semantics for a logic program with cut. The two semantics are based on different syntaxes, the first syntax (called *tree*) exploits a tree-like structure and is ideal both to have a graphical view of its evolution while the tree is being interpreted and to prove lemmas over it. The second syntax, called *elpi*, is the ELPI's syntax and has the advantage of reducing the computational cost of cutting and backtracking alternatives by using shared pointers. We aim to prove the equivalence of the two semantics together with some interesting lemmas of the cut behavior.

Notations used in the paper In the following section we note $\mathbb{B}$ for the boolean type, and $\top$ and $\perp$ for **true** and **false**. The option instances are noted $\bullet t$ for *Some t* and $\square$ for *None*.

se metti $r1 = g \ A \ B :- f \ A \ B$. allora $g \ e \ f$ sono funzioni, e puoi spiegare anche l'idea del detcheck qui

spiegare notazione for list

3 And-Or Tree semantics

The tree state The tree inductive, shown in Figure 4. We distinguish 5 main cases: KO , OK , and are special meta-symbols representing, respectively, the basic failed and a successful trees. These symbols are considered meta because they are internal intermediate symbols

```

Fixpoint get_end s A :  $\Sigma$  * tree :=
  match A with
  | TA _ | KO | OK  $\Rightarrow$  (s, A)
  | Or  $\square$   $\sigma_1$  B  $\Rightarrow$  get_end  $\sigma_1$  B
  | Or  $\bullet$ A _  $\Rightarrow$  get_end s A
  | And A _ B  $\Rightarrow$ 
    let (s', pA) := get_end s A in
    if pA == OK then get_end s' B
    else (s', pA)
  end.

Definition get_subst s A := (get_end s A).1.
Definition path_end A := (get_end  $\varepsilon$  A).2.
Definition success A := path_end A == OK.
Definition failed A := path_end A == KO.
Definition path_atom A :=
  if path_end A is TA _ then  $\top$ 
  else  $\perp$ .

```

(a) Definition of *get_end*(b) Functions from *get_end*

■ **Figure 5** *get_end* and derives functions

119 used to give structure to the tree.

120 The *TA* constructor (acronym for tree-atom) is the constructor of atoms in the tree. They
 121 are either the cut operator or a call, in these cases, the interpreter should evaluate the atom
 122 by either cutting the subtrees in the scope of the cut or backchaining the call in the program.

123 The two recursive cases of a tree are the *Or* and *And* nodes. The *Or* node $A \vee B_\sigma$ denotes
 124 a disjunction between two trees *A* and *B*. The first branch is optional, if absent it represents
 125 a dead tree, i.e. a tree that has been entirely explored. The second branch is annotated
 126 with a suspended substitution σ so that, upon backtracking to *B*, σ is used as the initial
 127 substitution for the execution of *B*.

128 The *And* node $A \wedge_{B_0} B$ represents a conjunction of two trees *A* and *B*. We call B_0 the
 129 reset point for *B*; it is used to restore the state of *B* to its initial form if a backtracking
 130 operation occurs on *A*. An example of this behavior is shown in Section 3.2

131 A graphical representation of a tree is shown in Figure 6a. To make the graph more
 132 compact, the *And* and *Or* nodes are n-ary rather than binary, and the children of these
 133 nodes associate to the right. The *KO* node acts as the neutral elements in the *Or* list, while
 134 *OK* is the neutral element of the *And* list.

135 *The tree interpreter* The interpretation of a tree is performed by two main routines: *step*
 136 and *next_alt* that traverse the tree depth-first, left-to-right. Then, then *runT* inductive
 137 makes the transitive closure of step *step* and *next_alt*: it iterates the calls to its these
 138 auxiliary functions as much as needed. In Types 7–9 we give the types contrats of these
 139 symbols where $\mathcal{F}_v$ is a set of variable names.


140 **Inductive** tag := Expanded | CutBrothers | Failed | Success. (6)

141 **Definition** step : $\mathbb{P} \rightarrow \mathcal{F}_v \rightarrow \Sigma \rightarrow \text{tree} \rightarrow (\mathcal{F}_v * \text{tag} * \text{tree}) :=$ (7)

142 **Definition** next_alt : $\mathbb{B} \rightarrow \text{tree} \rightarrow \text{option tree} :=$ (8)

143 **Inductive** runT (p : $\mathbb{P}$): $\mathcal{F}_v \rightarrow \Sigma \rightarrow \text{tree}$
 $\rightarrow \text{option } (\Sigma * \text{option tree}) \rightarrow \mathbb{B} \rightarrow \mathcal{F}_v \rightarrow \text{Prop} :=$ (9)

144 The tree interpreter, as in prolog, explores the tree in DFS strategy, to discover the
 145 substitution and the leaf of the tree that should be interpreted. These two pieces of
 146 information can be retrieved for a substitution *s* and a tree *A* thanks to the *get_end* routine,
 147 shown in Figure 5a If the tree is already a leaf, *get_end* returns its inputs. Otherwise, if the
 148 tree is a disjunction, the path continues on the left subtree, if it exists, otherwise *get_end*
 149 continues on the rhs using the substitution stored in the *Or* node. In the case of a conjunction,
 150 if *get_end* of the lhs returns the tree *OK*, then the *get_end* continues in the rhs, otherwise



■ **Figure 6** Some tree representations¹

151 we return the result of the lhs. We derive some useful definition from *get_end*. They are
 152 listed in Figure 5b.

153 We have pointed by a black arrow the *path_end* of all the trees in Figure 6.

154 3.1 The *step* procedure

155 The *step* procedure takes as input a program, a set of variables, a substitution, and a tree,
 156 and returns an updated set of variables, a *tag*, and an updated tree.

157 The set of variables v given in input to *step* are used and updated when a backchaining
 158 operation takes place. A call to *step* produces a good result if v contains all the variables
 159 in the tree. This guarantees that any call to *bc* returns a list of atoms whose variables are
 160 disjoint from the variables of the current working tree.

not true in fig-
 ure:fix + also fix
 in Section 2.1

¹ The substitution $\sigma_1 \dots \sigma_6$ are from Section 2.1, R_1, R_2 and R_3 are the reset points and correspond respectively to the list of atoms $[f \ Y \ Z, !], [!]$ and $[f \ Y \ Z]$

The *tag* (see Type 6) indicates the kind of an internal tree step: **CutBrothers** denotes the interpretation of a superficial cut, i.e., a cut whose parent nodes are all *And*-nodes. **Expanded** denotes the interpretation of non-superficial cuts or predicate calls. **Failure** and **Success** are returned for, respectively, *failed* and *success* trees.

The step procedure is intended to evaluate atoms, that is, it leaves unchanged the tree iff its *path_end* is not an atom, i.e. if it is a success- or a failed-tree.

Lemma succ_step_iff: $\forall p \ v \ s \ A, \text{success } A \leftrightarrow \text{step } p \ v \ s \ A = (v, \text{Success}, A).$ (1)

Lemma fail_step_iff: $\forall p \ v \ s \ A, \text{failed } A \leftrightarrow \text{step } p \ v \ s \ A = (v, \text{Failed}, A).$ (2)

The interpretation of a call to a term *c* starts by calling the *bc* function on *c*. The result of the backchain is a list *l*. If *l* is empty then the current call is replaced by the *KO* node, otherwise it is replaced by a right skewed tree made by *Or* inner-nodes and each leave correspond to a list of atom *e* in *l*. If *e* is empty then it is mapped to the *OK* tree, otherwise it is mapped to a right-skewed tree made by *And* inner-nodes.

In Figure 6a, we provide a clear example of a call to *step* on the query $q \ 0 \ R$ on the program *P* in Figure 3. Let *c*₀, *c*₁ and *c*₂ be the children of the root. *c*₀, by construction is the *None* tree, while *c*₁ and *c*₂ correspond respectively to the premises of rules *r*₁ and *r*₂ in *P*. We need the first *None* subtree so that we can attach a substitution to *c*₁. In particular, we note that the *c*₁ (resp. *c*₂) are related to the substitution σ_1 (resp. σ_2). The value of σ_1 and σ_2 are the same as in Section 2.1. The reset points in the tree are marked with the *R* symbol. $R_1 := [f \ Y \ Z, !]$, $R_2 := [!]$ and $R_3 := f \ Y \ Z$, they are essentially the list of atoms that allowed to generate the current subtree. Other example of call to the *step* procedure triggering a backchain are Figures 6b, 6c, and 6e.

The interpretation of a cut has three main impacts: at first the cut is replaced by the *OK* node, then some special subtrees, in the scope of the *Cut*, are cut away: in particular we need to soft-kill the left-siblings of the *Cut* and hard-kill the right-uncles of the *Cut*.

► **Definition 1** (Left-siblings (resp. right-sibling) and Right-uncles). *Given a node A, the left-siblings (resp. right-sibling) of A are the list of subtrees sharing the same parent of A and that appear on its left (resp. right). The right-uncles of A are the list of right-sibling of the father of A.*

► **Definition 2** (Hard-kill, $!_r$ and Soft-kill, $!_l$). *Given a tree t, hard-kill replaces the given subtree with the KO node. If t is a success-tree, then soft-kill replaces with KO all subtrees that are not part of the path in t leading to the OK node.*

An example of the impact of the cut is show in Figure 6f. In Figure 6e, we see that the *path_end* is a cut. The cut is replaced by the *OK* node and it need to cut-away some alternatives in the tree. The red arrows indicate the scope of the cut, they are labeled with $!_r$ and $!_l$ to explain which kind of cut has been performed on each pointed subtree. We note that, after the evaluation of a cut, the path leading to the cut node is not affected by $!_l$, the success under the substitution σ_4 is kept, while the success under the substitution σ_5 is replaced by the *KO* node. If we refer to the paragraph in Section 2.1, we are discarding the third rule of the *r* predicate.

3.2 The next_alt procedure

The *step* procedure alone is not sufficient to reproduce entirely the behavior of the full expected prolog interpreter. As proved in Lemma 1 and Lemma 2, *step* does not make any progression on the tree being interpreted. To solve this problem, we need a routine capable

sottolineare
sostituzioni

of backtracking in case of failures. This routine return the backtracked tree if it exists, *None* otherwise.

In Figure 6c, we are facing a failure: the interpreter has evaluated the atom $\mathbf{f} \ X \ Y$ under the substitution σ_3 which assigns Y to 0, but, in Figure 3, we have no rule satisfying this query. The role of *next_alt*, in this example, is to kill the path leading to this failure, while resetting the failure to its original shape. We see how the reset point plays a major role in this case. At first we are killing the *OK* node under the substitution σ_3 , since this tree have been entirely explored and led to no solution, then we replace the *KO* node with the information stored in R_1 . That is, thanks to R_1 , we are restoring the $\mathbf{f} \ Y \ Z$ call, so that we are able to evaluate this node under the substitution σ_4 . More concretely, we have uncessesfully tried rule $\mathbf{r} \ 0 \ 0$, therefore, we need to continue by trying rule $\mathbf{r} \ 0 \ 1$.

Furthermore, *next_alt* accomplishes to another important task. Its behavior changes depending to the value of the boolean it receives in entry: *next_alt false A* aims to recursively cancel all the failures (if any) in A , whereas, *next_alt true A* aims to cancel the first success in A . This second behavior is helpful since, we need to progress the tree in case of success. For example, consider the tree in Figure 6f, it contains a success, mapping namely R to 2. The interpreter should return this sucess, but, due to the “non-deterministic” behavior of the prolog system, we also need to return all the other solution that could be available by evaluating the rest of the tree. After the replacement of the *path_end* in Figure 6f with *KO*, we remain with a tree having all failing path. The *next_alt* procedure keeps removing them until discovering that the tree have no worth-to explore paths, therefore it will return *None*.

Some interesting property of *next_alt* are shown below and allow to see how *next_alt* complements *step* wrt failed, success and *path_atom* trees.

Lemma *path_atom_next_alt_id*: $\forall b \ A, \text{path_atom } A \rightarrow \text{next_alt } b \ A = \bullet A.$ (3)

Lemma *next_alt_failedF*: $\forall b \ A \ A', \text{next_alt } b \ A = \bullet A' \rightarrow \text{failed } A' = \perp.$ (4)

Lemma *next_altFN_fail*: $\forall A, \text{next_alt } \perp \ A = \square \rightarrow \text{failed } A.$ (5)

3.3 The *runT* inductive

The inductive *runT* is modeled as a function taking 4 inputs and returning 3 outputs. The inputs are the program, a set of variables v_0 , an initial substitution σ_0 , and a tree t_0 . The outputs are an optional object r containing a substitution σ_1 and an optional tree t_1 , and a boolean b and a set of updated variables v_1 . The very interesting piece of information is the result r , if *None*, it tells that the interpretation failed, otherwise, σ_1 represents the most-general unificator that makes the execution of the tree t_0 under the initial substitution σ_0 succeeds: σ_1 is an extension of σ_0 . The tree t_1 is the updated tree containing the alternatives that have not yet been explored, if these alternatives do not exists, t_1 is *None*. The output b and v_1 are there to tell respectively if during the execution a superficial cut has been evaluated, and what is the updated set of variables after the interpretation.

The procedure *runT* is based on four main derivation rules, shown in Figure 7. (STOPT) If the *path_end* of the tree t_0 is a success, the substitution σ_1 is retrived in t_0 from σ_0 and the input tree is cleaned of its successful path. (STEPT) If the *path_end* of t_0 is an atom, then *step* is invoked to evaluate t_0 , and *runT* is recursively called on the new tree. (BACKT) If the *path_end* of t_0 is a failure, *next_alt* is called to clear the failed path; if the resulting cleaned tree t_1 exists, *runT* is recursively called on it. Finally, (FAILT) takes into account tree with no solutions.

Lemma *runT_det*: $\forall p \ v_0 \ \sigma_0 \ t_0 \ r_1 \ r_2 \ b_1 \ b_2 \ v_1 \ v_2,$
 $\text{runT } p \ v_0 \ \sigma_0 \ t_0 \ r_1 \ b_1 \ v_1 \rightarrow \text{runT } p \ v_0 \ \sigma_0 \ t_0 \ r_2 \ b_2 \ v_2 \rightarrow r_2 = r_1 \wedge v_2 = v_1 \wedge b_2 = b_1.$ (6)

250 The lemma above tells that $runT$ behaves deterministically. The proof is by induction
 251 on the first $runT$ execution and the proof is rather immediate, since all rules are in exclusive
 252 due to the hypothesis on the $path_end$: the hypotheses $success\ t_0$, $path_atom\ t_0$ and $failed\ t_0$
 253 are exclusive. Moreover, by Lemma 6, we have that the hypothesis of **FAILT** implies $failed\ t_0$,
 254 therefore it is exclusive wrt **STOPT** and **STEPT**, but also from **BACKT** since they have
 255 different expectations on the output of $next_alt$.

$$\begin{array}{c}
 \frac{success\ t_0 \quad get_subst\ \sigma_0\ t_0 = \sigma_1 \quad next_alt\ \top\ t_0 = t_1}{runT\ v_0\ \sigma_0\ t_0\ \bullet(\sigma_1, t_1) \perp v_0} \text{STOPT} \\
 \\
 \frac{path_atom\ t_0 \quad b_2 = is_cb\ tg \parallel b_1 \quad step\ p\ v_0\ \sigma_0\ t_0 = (v_1, tg, t_1) \quad runT\ v_1\ \sigma_0\ t_1\ r\ b_1\ v_2}{runT\ v_0\ \sigma_0\ t_0\ r\ b_2\ v_2} \text{STEPT} \\
 \\
 \frac{failed\ t_0 \quad next_alt\ \perp\ t_0 = \bullet t_1 \quad runT\ v_0\ \sigma_0\ t_1\ r\ n\ v_1}{runT\ v_0\ \sigma_0\ t_0\ r\ n\ v_1} \text{BACKT} \\
 \\
 \frac{next_alt\ \perp\ t_0 = \square}{runT\ v_0\ \sigma_0\ t_0\ \square \perp v_0} \text{FAILT}
 \end{array}$$

■ **Figure 7** Rule system for $runT$

256 3.4 Properties of $runT$

257 Working with $runT$ allows to prove some interesting properties about the behavior of the
 258 cut operator wrt disjunction, i.e. the Or node. We have proven several properties that are
 259 available at [this](#) link. Below we show two of these lemmas.

260 **Lemma runSST_or:** $\forall\ p\ v_0\ v_1\ \sigma_0\ \sigma_1\ t_0\ t'_0\ sm\ t_1,$
 $runT\ p\ v_0\ \sigma_0\ t_0\ \bullet(\sigma_1, \bullet t'_0) \top v_1 \rightarrow$
 $runT\ p\ v_0\ \sigma_0\ (Or\ \bullet t_0\ sm\ t_1)\ \bullet(\sigma_1, \bullet(Or\ \bullet t'_0\ sm\ KO)) \perp v_1.$ (7)

261 **Lemma runNT_or:** $\forall\ p\ v_0\ v_1\ \sigma_0\ t_0\ sm\ t_1,$
 $runT\ p\ v_0\ \sigma_0\ t_0\ \square \top v_1 \rightarrow$
 $runT\ p\ v_0\ \sigma_0\ (Or\ \bullet t_0\ sm\ t_1)\ \square \perp v_1.$ (8)

262 **Lemma run_orSST** $p\ v_0\ v_2\ \sigma_0\ \sigma_1\ sm\ t_0\ t'_0\ t_1\ t'_1:$
 $runT\ p\ v_0\ \sigma_0\ (Or\ \bullet t_0\ sm\ t_1)\ \bullet(\sigma_1, \bullet(Or\ \bullet t'_0\ sm\ t'_1)) \perp v_2 \rightarrow$
 $\exists\ b, runT\ p\ v_0\ \sigma_0\ t_0\ \bullet(\sigma_1, \bullet t'_0)\ b\ v_2 \wedge t'_1 = \text{if } b \text{ then } KO \text{ else } t_1.$ (9)

263 In Lemma 7, we evaluate the tree t_0 . This evaluation produces a success with signature σ_1
 264 and a new tree t'_0 . During the evaluation, a superficial cut in t_1 is traversed. This implies that
 265 the evaluation of $\bullet t_0 \vee_{sm} t_1$ returns σ_1 as substitution, and the alternatives are $\bullet t'_0 \vee_{sm} KO$.
 266 The right-hand side of the Or node is replaced with KO by $!_r$.

267 In Lemma 8, we are in a similar scenario: t_0 evaluates a superficial cut but, this time,
 268 produces a failure. This implies that the evaluation of $\bullet t_0 \vee_{sm} t_1$ also fails.

269 Finally, in Lemma 9, we evaluate $\bullet t_0 \vee_{sm} t_1$. The evaluation succeeds, producing the
 270 result $\bullet(\sigma_1, \bullet t'_0 \vee_{sm} t'_1)$. This means that t_0 still has t'_0 as unexplored alternatives. We deduce
 271 that the evaluation of t_0 produces the substitution σ_1 with alternatives t'_0 . However, we have
 272 no information about whether a superficial cut has been crossed in t_0 . If this is the case, we
 273 know that t'_1 is cut away by $!_r$; otherwise, t'_1 has not been explored and is equal to t_1 .

Definition $\text{stepE } v_0 \text{ t s a gl} :=$
 $\text{let } (v_1, \text{rs}) := \text{bc p } v_0 \text{ t s in}$
 $(v_1, \text{save_alts a gl rs}).$

$$\frac{}{\text{runE } v_0 ((\sigma_0, \emptyset) :: a) \bullet (\sigma_0, a)} \text{STOPE}$$

$$\frac{\text{runE } v_0 ((\sigma_0, g) :: \text{ca}) \text{ r}}{\text{runE } v_0 ((\sigma_0, (\text{cut}, \text{ca}) :: g) :: _) \text{ r}} \text{CUTE}$$

$$\frac{\text{stepE } v_0 \text{ t } \sigma_0 \text{ a g} = (v_1, \text{bs}) \quad \text{runE } v_1 (\text{bs} ++ a) \text{ r}}{\text{runE } v_0 ((\sigma_0, (\text{call t}, _) :: g) :: a) \text{ r}} \text{CALLE}$$

$$\frac{}{\text{runE } _ \emptyset \square} \text{FAILE}$$

■ **Figure 8** Rule system for runE

274 From these lemmas, it is clear how the order in which alternatives are evaluated changes
 275 the behavior of the program.

276 4 Stack semantics

277 We now introduce the ELPI semantics. The interpreter is close to that of the ELPI language
 278 and provides an operational semantics that is similar to the one proposed by Pusch in [16],
 279 which is in turn closely related to the semantics given by Debray and Mishra in [6, Section 4.3].
 280 Pusch mechanized this semantics in Isabelle/HOL, together with several optimizations that
 281 are also present in the Warren Abstract Machine [20, 1].

282 We represent a state of the ELPI language thanks to the two mutual inductives shown
 283 below.

Inductive $\text{alts} := \text{nilA} \mid \text{consA of } (\Sigma * \text{goals}) \ \& \ \text{alts}$
with $\text{goals} := \text{nilG} \mid \text{consG of } (A * \text{alts}) \ \& \ \text{goals} .$

284 An ELPI state is an enhanced two-dimensional list. The outermost list represents a list
 285 of alternatives in disjunction, each accompanied by the substitution that should be used
 286 for their interpretation. The innermost list is a list of atoms, that is, the list of goals in
 287 conjunction. These goals are decorated with a pointer to an alternative list, which is used to
 288 keep track of where to jump when a cut is interpreted. We call these special alternatives the
 289 *cut-to* alternatives.

290 The ELPI interpreter is called runE and has the following type.

291 **Inductive** $\text{runE} : \mathcal{F}_v \rightarrow \text{alts} \rightarrow \text{option } (\Sigma * \text{alts}) \rightarrow \text{Prop} :=$ (10)

292 The inductive runE represents a routine taking as input a set of variables and a list
 293 of alternatives and returning an option. $\square$ stands for interpretation with no solution,
 294 whereas $\bullet(\sigma, a)$ represent a successful interpretation with σ as solution and a as the list of
 295 non-yet-explored choice points.

296 The rule system for runE is given in Figure 8 and is made by 4 cases. We distinguish
 297 two main cases: if the list of alternatives is empty, (FAILE) we have a failure: we stop the
 298 computation and return $\square$. If the list of alternatives is the list $(\sigma_0, h) :: a$, the interpreter
 299 evaluates the first choice point h under the substitution σ_0 . If h is empty, (STOPE), then we

se riesco success
 (A or B) = suc-
 cess (B or A) if
 no sup cut in A
 and B

300 have a success, i.e. we return the input σ_0 and leave a as the future choice points. Otherwise,
 301 h is the list $(t, ca) :: g$ where t is the first atom to evaluate, ca is its cut-alternative list and g
 302 is the list of future conjuncts. If t is a cut (CUTE), we need to keep evaluating g under σ_0
 303 but we change the alternatives to ca . Finally, in the case of a call (CALLE), we backchain
 304 the call in the program. If rs is the list of rule premises returned by `bc`, then each list of
 305 premises if mapped to a list of goals (`goals`) have a as cut-alternative and the global list is
 306 translated to a list of alternative (`alts`).

307 *Example of execution*

308 In the example only cuts have pointer. Call do not use them

309 *Remark* It is not possible to prove the same lemma as in section xxx, because no structure,
 310 the scope of the cut ... spiegare bene il problema so that i can reuse this subsection has a
 311 hook when I will say that l2t is difficult + Yves

312 5 Equivalence of the two semantics

313 The equivalence between the two semantics is possible under two conditions: we need to
 314 work with “valid tree”, i.e. all of those trees that can be generated from a call. Secondly,
 315 we need to translate trees into elpi states. In the next to subsection we propose the two
 316 functions *valid_state* and *t2l*.

317 5.1 Valid trees

318 The inductive tree allows one to generate a large number of trees, some of which are not
 319 valid, in the sense that they cannot be produced starting from a given query. The class of
 320 valid trees is characterized by the function shown in Figure 9.

321 While all base cases of tree are considered valid, we need to analyze carefully the cases
 322 for the *Or* and *And* constructors.

323 For the *Or* constructor, we distinguish two cases depending on whether the left-hand
 324 side (lhs) exists. If it does not exist, then the right-hand side (rhs) must be a valid tree.
 325 Otherwise, the lhs must itself be a valid tree, and the rhs is either the *KO* tree, since it may
 326 have been removed by the evaluation of a superficial cut in the lhs, or it has not yet been
 327 explored. In the latter case, it is a `base_or` tree, namely the right-skewed tree formed by a
 328 disjunction of conjunctions that is generated after a successful backchain to a call.

329 For the *And* constructor, the lhs is required to be a valid tree. The shape of the rhs
 330 depends on whether the lhs is a success. If the lhs is not successful, then the rhs has never
 331 been explored: the procedures *step* and *next_alt* modify the rhs only when the lhs succeeds.
 332 In this case, the lhs must be the right-skewed tree containing conjunctions of premises atoms,
 333 the reset point B_0 ensure what shape the rhs should have. If the lhs is a success tree, then
 334 the rhs can be modified by *step* and *next_alt*, therefore it must be a valid tree.

335 5.2 From trees to lists

336 The translation of a tree to a list is shown in ???. It takes the tree to be translated, a
 337 substitution (called s), a list of alternatives, called bt . The substitution s tells what is the
 338 substitution of the alternative we are building and is updated when going to the rhs of the
 339 *Or* constructor. The bt list represents the future alternatives list that have the same depth
 340 of the current tree root. They are useful to know if they should be added to the current goal
 341 as cut alternatives.

```

Fixpoint valid_tree A :=
  match A with
  | TA _ | OK | KO => T
  | Or □ _ B => valid_tree B
  | Or •A _ B => valid_tree A &&
    ((B == KO) || B.base_or B)
  | And A B0 B => valid_tree A &&
    if success A then valid_tree B
    else B == big_and B0
  end.

```

■ **Figure 9** Valid tree

342 An *OK* node represent a success in a tree, then the translation of this tree returns a
 343 singleton list with the cuple $(s, [::])$: there is one alternative with 0 goals, i.e. a success in
 344 elpi by STOPE. The *KO* node represent a failure, and, therefore 0 alternatives. An atom is
 345 translated into a singleton with the atom as first goal and the empty cut-alternative list: we
 346 are not adding *bt* since they are the alternatives for the right-sibling, not the right-uncles,
 347 this means that if we have *a* is a cut, then it erases the alternatives *bt*.

348 In a disjunction, we are scendendo di un livello the distance from our backtracking list.
 349 This means that *bt* becomes the list of right-uncles of the two branches of the *Or* constructor.
 350 The compilation of the rhs is done independently of if the lhs exists: we transform the rhs
 351 into an elpi state and passing the empty list of alternatives, since the rhs as no right-siblings.
 352 Then, thanks to `add_ca_deep`, we recursively add *bt* to every cut-alternative appearing in
 353 the translated goals. If the lhs exists, we translate it and pass the translation of the rhs as
 354 the list of right-siblings. Then we prepend the translated lhs to the translated rhs and add
 355 *bt* with `add_ca_deep`.

356 We compile the *And* case by translating its lhs to the list *lA* and reset point to the list
 357 *lB0*. Then, thanks to `add_deep`, we recursively append *lB0* to the alternatives born in *lA*,
 358 i.e. we leave unchanged the pointers to *bt*, if any, in *lA*. We finally need to treat cases
 359 whenever the list *lA* is empty or not. In the former case the empty list is returned, otherwise,
 360 we have a list $[:: (s_0, gs) \& al]$. By the semantics of the *And* in the *runT* interpreter, the rhs
 361 of the *And* represent the sequent of goals to be executed after the first alternative in the
 362 lhs, it corresponds to the list of goals *gs*. The reset point *B0* is to be appended to the next
 363 alternatives in the lhs, that is the queue of *lA*, that is *al*.

alternative in
tree? rephrase

364 We note therefore that the compilation of *lB* is done by passing the alternatives *al* + *bt*
 365 since the alternatives in *al* are to be...

366 5.3 Equivalence theorems


367 We have now all the ingredients that ...

```

Lemma tree_to_elpi fv A σ1 b fv' r :
  vars_tree A `<=>` fv → vars_sigma σ1 `<=>` fv →
  valid_tree A →
  runT u p fv σ1 A r b fv' →
  let xs := t2l A σ1 ∅ in
  let r' := omap (fun '(s, t) => (s, (t2l (odflt KO t) σ1 ∅))) r in
  runE u p fv xs r'.

```

(10)



■ **Figure 10** Induction scheme for Lemma 11

```

Lemma elpi_to_tree v0 a r :
  runE u p v0 a r →
  ∀ σ0 t, valid_tree t → t2l t σ0 ∅ = a →
  ∃ b v1,
  if r is •(σ1, a') then
    ∃ t1, runT u p v0 σ0 t •(σ1, t1) b v1 ∧ t2l (odflt KO t1) σ0 ∅ = a'
  else runT u p v0 σ0 t □ b v1.

```

(11)

The proof of Lemma 11 is based on the idea explained in [2, Section 3.3]. An ideal statement for this lemma would be: given a function 12t transforming an elpi state to a tree, we would have that the execution of an elpi state e is the same as executing runT on the tree resulting from $\text{12t}(e)$. However, it is difficult to retrieve the structure of an elpi state and create a tree from it. This is because, in an elpi state, we have no clear information about the scope of an atom inside the list and, therefore, no evident clue about where this atom should be placed in the tree.

Our theorem states that, starting from a valid tree t which translates to a list of alternatives $(\sigma_1, g) :: a$. If we run in elpi the list of alternatives, then the execution of the tree t returns the same result as the execution in elpi. The proof is performed by induction on the derivations of the elpi execution. We have 4 derivations.

We have 4 cases to analyse:

6 Case study: determinacy analysis

we mechanize the first order part of xxx.

snippet det, main thm, invariant det tree (valid tree prev section?)

proof induction on exec, step/next alt preserving invariant proved by induction on the tree.

with list semantics cut and next alt requires to express a link between the current alt and the current goal, which is non trivial without an intermediate data structure like the tree

7 Related work

The only mechanization of Prolog's semantics we could find is the one by Pusch in ISABELLE/HOL [16] that is based on the model of Debray and Mishra [6, Sec. 4.3]. Their work starts from that semantics and refines it by introducing some optimizations usually found in the Warren abstract machine [20, 1]. They do not use the semantics to prove properties on the execution of programs as we do in our use case, hence they do not face the difficulty described in section XXX that pushed us to develop a more explicit semantics (with respect to cut). In some way our work is complementary, showing the stack based semantics equivalent to a more optimized one, we could consider implementing it in Elpi.

Another relevant work are the ones by King cite, but we could not find their code. Their semantics is denotational and cannot easily cover all the features our paper XX does, but would have been sufficient for this first step.

The proof technique we use to relate the two semantics is inspired by XXX and we thank Y.Bertot for insightful discussions on the topic. In XXX Bertot relates bla bla, witout never constructing a decompiler, exactly as we never build a list to tree function.

8 Conclusion

We do this. This is a first step for XXXX. We need to add substitutions. The missing part of the proof becomes related to abstract interpretation, since the det types for a lattice and one has to prove that the concrete runtime values variables take agree with the abstraction computed statically. This proof is ongoing and somehow orthogonal to the current one since the cut operator and the HO feature do not interfere with eachother in complex ways.

References

- 1 Hassan Aït-Kaci. *Warren's Abstract Machine: A Tutorial Reconstruction*. The MIT Press, 08 1991. doi:10.7551/mitpress/7160.001.0001.
- 2 Yves Bertot. A certified compiler for an imperative language. Technical Report RR-3488, INRIA, September 1998. URL: <https://inria.hal.science/inria-00073199v1>.
- 3 Valentin Blot, Denis Cousineau, Enzo Crance, Louise Dubois de Prisque, Chantal Keller, Assia Mahboubi, and Pierre Vial. Compositional pre-processing for automated reasoning in dependent type theory. In Robbert Krebbers, Dmitriy Traytel, Brigitte Pientka, and Steve Zdancewic, editors, *Proceedings of the 12th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2023, Boston, MA, USA, January 16-17, 2023*, pages 63–77. ACM, 2023. doi:10.1145/3573105.3575676.
- 4 Cyril Cohen, Enzo Crance, and Assia Mahboubi. Trocq: Proof transfer for free, with or without univalence. In Stephanie Weirich, editor, *Programming Languages and Systems*, pages 239–268, Cham, 2024. Springer Nature Switzerland.
- 5 Cyril Cohen, Kazuhiko Sakaguchi, and Enrico Tassi. Hierarchy Builder: Algebraic hierarchies Made Easy in Coq with Elpi. In *Proceedings of FSCD*, volume 167 of *LIPIcs*, pages 34:1–34:21, 2020. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.FSCD.2020.34>, doi:10.4230/LIPIcs.FSCD.2020.34.
- 6 Saumya K. Debray and Prateek Mishra. Denotational and operational semantics for prolog. *J. Log. Program.*, 5(1):61–91, March 1988. doi:10.1016/0743-1066(88)90007-6.
- 7 Cvetan Dunchev, Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. ELPI: fast, embeddable, λProlog interpreter. In *Proceedings of LPAR*, volume 9450 of *LNCS*, pages 460–468. Springer, 2015. URL: <https://inria.hal.science/hal-01176856v1>, doi:10.1007/978-3-662-48899-7_32.
- 8 Davide Fissore and Enrico Tassi. A new Type-Class solver for Coq in Elpi. In *The Coq Workshop*, July 2023. URL: <https://inria.hal.science/hal-04467855>.
- 9 Davide Fissore and Enrico Tassi. Higher-order unification for free!: Reusing the meta-language unification for the object language. In *Proceedings of PPDP*, pages 1–13. ACM, 2024. doi:10.1145/3678232.3678233.
- 10 Davide Fissore and Enrico Tassi. Determinacy checking for elpi: an higher-order logic programming language with cut. In *Practical Aspects of Declarative Languages: 28th International Symposium, PADL 2026, Rennes, France, January 12–13, 2026, Proceedings*, pages 77–95, Berlin, Heidelberg, 2026. Springer-Verlag. doi:10.1007/978-3-032-15981-6_5.
- 11 Benjamin Grégoire, Jean-Christophe Léchenet, and Enrico Tassi. Practical and sound equality tests, automatically. In *Proceedings of CPP*, page 167–181. Association for Computing Machinery, 2023. doi:10.1145/3573105.3575683.

- 446 **12** Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. Implementing type theory
447 in higher order constraint logic programming. In *Mathematical Structures in Computer*
448 *Science*, volume 29, pages 1125–1150. Cambridge University Press, 2019. doi:10.1017/
449 S0960129518000427.
- 450 **13** Robbert Krebbers, Luko van der Maas, and Enrico Tassi. Inductive Predicates via Least
451 Fixpoints in Higher-Order Separation Logic. In Yannick Forster and Chantal Keller, editors,
452 *16th International Conference on Interactive Theorem Proving (ITP 2025)*, volume 352 of *Leib-*
453 *niz International Proceedings in Informatics (LIPIcs)*, pages 27:1–27:21, Dagstuhl, Germany,
454 2025. Schloss Dagstuhl – Leibniz-Zentrum für Informatik. URL: [https://drops.dagstuhl.](https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.ITP.2025.27)
455 [de/entities/document/10.4230/LIPIcs.ITP.2025.27](https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.ITP.2025.27), doi:10.4230/LIPIcs.ITP.2025.27.
- 456 **14** Dale Miller. A logic programming language with lambda-abstraction, function variables, and
457 simple unification. In *Extensions of Logic Programming*, pages 253–281. Springer, 1991.
- 458 **15** Dale Miller and Gopalan Nadathur. *Programming with Higher-Order Logic*. Cambridge
459 University Press, 2012.
- 460 **16** Cornelia Pusch. Verification of compiler correctness for the wam. In Gerhard Goos, Juris
461 Hartmanis, Jan van Leeuwen, Joakim von Wright, Jim Grundy, and John Harrison, editors,
462 *Theorem Proving in Higher Order Logics*, pages 347–361, Berlin, Heidelberg, 1996. Springer
463 Berlin Heidelberg.
- 464 **17** Enrico Tassi. Elpi: an extension language for Coq (Metaprogramming Coq in the Elpi λ Prolog
465 dialect). In *The Fourth International Workshop on Coq for Programming Languages*, January
466 2018. URL: <https://inria.hal.science/hal-01637063>.
- 467 **18** Enrico Tassi. Deriving proved equality tests in Coq-Elpi. In *Proceedings of ITP*, volume 141 of
468 *LIPIcs*, pages 29:1–29:18, September 2019. URL: <https://inria.hal.science/hal-01897468>,
469 doi:10.4230/LIPIcs.CVIT.2016.23.
- 470 **19** Luko van der Maas. Extending the Iris Proof Mode with inductive predicates using Elpi.
471 Master’s thesis, Radboud University Nijmegen, 2024. doi:10.5281/zenodo.12568604.
- 472 **20** David H.D. Warren. An Abstract Prolog Instruction Set. Technical Report Technical Note 309,
473 SRI International, Artificial Intelligence Center, Computer Science and Technology Division,
474 Menlo Park, CA, USA, October 1983. URL: [https://www.sri.com/wp-content/uploads/](https://www.sri.com/wp-content/uploads/2021/12/641.pdf)
475 [2021/12/641.pdf](https://www.sri.com/wp-content/uploads/2021/12/641.pdf).