

# **Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis**

**Jane Open Access**  

Dummy University Computing Laboratory, [optional: Address], Country

My second affiliation, Country

**Joan R. Public<sup>1</sup>**  

Department of Informatics, Dummy College, [optional: Address], Country

---

## **Abstract**

9 Lorem ipsum dolor sit amet, consectetur adipiscing elit. Praesent convallis orci arcu, eu mollis dolor.  
10 Aliquam eleifend suscipit lacinia. Maecenas quam mi, porta ut lacinia sed, convallis ac dui. Lorem  
11 ipsum dolor sit amet, consectetur adipiscing elit. Suspendisse potenti.

12 **2012 ACM Subject Classification** Replace ccsdesc macro with valid one

13 **Keywords and phrases** Dummy keyword

14 **Digital Object Identifier** 10.4230/LIPIcs.CVIT.2016.23

15 **Funding Jane Open Access:** (Optional) author-specific funding acknowledgements

16 **Joan R. Public:** [funding]

17 **Acknowledgements** I want to thank ...

## **1 Introduction**

19 ELPI is a dialect of  $\lambda$ PROLOG (see [14, 15, 7, 12]) used as an extension language for the ROCQ  
20 prover (formerly the Coq proof assistant). ELPI has become an important infrastructure  
21 component: several projects and libraries depend on it [13, 3, 4, 19, 8, 9]. Examples include  
22 the Hierarchy-Builder library-structuring tool [5] and Derive [17, 18, 11], a program-and-proof  
23 synthesis framework with industrial applications at SkyLabs AI.

24 Starting with version 3, ELPI gained a static analysis for determinacy [10] to help users  
25 tame backtracking. ROCQ users are familiar with functional programming but not necessarily  
26 with logic programming and uncontrolled backtracking is a common source of inefficiency  
27 and makes debugging harder. The determinacy checkers identifies predicates that behave  
28 like functions, i.e., predicates that commit to their first solution and leave no *choice points*  
29 (places where backtracking could resume).

30 This paper reports our first steps towards a mechanization, in the ROCQ prover, of the  
31 determinacy analysis from [10]. We focus on the control operator *cut*, which is useful to  
32 restrict backtracking but makes the semantic depart from a pure logical reading.

33 We formalize two operational semantics for PROLOG with cut. The first is a stack-  
34 based semantics that closely models ELPI's implementation and is similar to the semantics  
35 mechanized by Pusch in ISABELLE/HOL [16] and to the model of Debray and Mishra [6,  
36 Sec. 4.3]. This stack-based semantics is a good starting point to study further optimizations  
37 used by standard PROLOG abstract machines [20, 1], but it makes reasoning about the scope  
38 of *cut* difficult. To address that limitation we introduce a tree-based semantics in which the  
39 branches pruned by *cut* are explicit and we prove the two semantics equivalent. Using the

---

<sup>1</sup> Optional footnote, e.g. to mark corresponding author



© Jane Open Access and Joan R. Public;  
licensed under Creative Commons License CC-BY 4.0

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:10

 Leibniz International Proceedings in Informatics

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

## 23:2 Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

```

Inductive P := IP of nat. Inductive D := ID of nat. Inductive V := IV of nat.

Inductive Tm :=
| Tm_P of P      | Tm_D   of D      | Tm_V of V      | Tm_App of Tm & Tm.

Inductive Callable :=
| Callable_P of P | Callable_App of Callable & Tm.

```

Figure 1 Tm and Callable types

40 tree-based semantics we then show that if every rule of a predicate passes the determinacy  
 41 analysis, the call to a deterministic predicate does not leave any choice points.

## 2 Common code: the language

put unif and progs  
 gram in variables  
 hides from types  
 46 the them. The smallest unit of code that we can use in the langauge is an atom. The atom  
 inductive (see Type 1) is either a cut or a call. A call carries a callable term (see Figure 1).  
 47 A term (Tm) is either a predicate, a datum, a variable or the binary application of a term to  
 another. A Callable is a term accepting predicates only predicates as functors.

```

48 Inductive A := cut | call : Callable -> A.                               (1)
49 Record R := mkR { head : Callable; premises : list A }.                  (2)
50 Record program := { rules : seq R; sig : sigT }.                         (3)
51 Definition Sigma := {fmap V -> Tm}.                                       (4)
52 Definition bc : Unif -> program -> fvS -> Callable ->
  Sigma -> fvS * seq (Sigma * R) :=                                         (5)

```

53 A rule (see Type 2) is made a head of type term and a list of premises, the premises are  
 54 atoms. A program (see Type 3) is made by a list of rules and a mapping from predicates to  
 55 their signatures. The type sigT is the classic type from the simply typed lambda calculus, i.e.  
 56 it is either a base type or an arrow. We decorate arrows to know the mode of the lhs type.

57 A substitution (see Type 4) is a mapping from variables to terms. It is the output of a  
 58 successful query and is often called the output of a query.

```

Record Unif := {
  unify : Tm -> Tm -> Sigma -> option Sigma;
  matching : Tm -> Tm -> Sigma -> option Sigma;
}.

```

59 The backchain function (bc, see Type 5) filters the rules in the program that can be  
 60 used on a given query. It takes: a unificator  $U$  which explains how to unify terms up to  
 61 standard unification (for output terms) or matching (for input terms); a program  $P$  to explore  
 62 and filter; a set  $S$  of free variable (fvS) allowing to fresh the program  $P$  by renaming the  
 63 its variables; a query  $q$ ; and the substitution  $\sigma$  in which the query  $q$  lives. The result of a  
 64 backchain operation is couple made of an extension of  $S$  containing the new variables that  
 65 have been allocated during the unification phase and a list of filtered rules  $r$  accompagnate  
 66 by their a substition. This substitution is the result of the unification of  $q$  with the head of  
 67 each rule in  $r$ .

68 In Figure 2, we have an example of a simple ELPI program which will be used in the  
 69 following section of the paper as an example to show how backtracking and the cut operator  
 70 works in the semantcis we propose. The translation of these rules in the ROCQ representation  
 71 is straightforward.

```
f 1 2.    f 2 3.    r 2 4.    r 2 8.
g X X.          % r1
g X Z :- r X Z, !.   % r2
g X Z :- f X Y, f Y Z.   % r3
```

Figure 2 Small ELPI program example

## 2.1 The cut operator

The semantics of the cut operator adopted in the ELPI language corresponds to the *hard cut* operator of standard SWI-PROLOG. This operator has two primary purposes. First, it eliminates all alternatives that are created either simultaneously with, or after, the introduction of the cut into the execution state.

To illustrate this high-level description, consider the program shown in Figure 2 and the query  $q = g 2 \text{ Z}$ . All three rules for  $g$  can be used on the query  $q$ . They are tried according to their order of appearance in the program: rule  $r_1$  is tried first, followed by  $r_2$ , and  $r_3$ .

The first rule has no premises and immediately returns the assignment  $\text{Z} = 2$ . However, the computation does not terminate at this point, since two additional unexplored alternatives remain, corresponding to the premises of rules  $r_2$  and  $r_3$ .

The premises of rule  $r_2$  are  $\text{r } 2 \text{ Z}, !$ . At this stage, the role of the cut becomes apparent. If the premise  $\text{r } 2 \text{ Z}$  succeeds, the cut commits to this choice and removes the premises of rule  $r_3$  from the alternative list, as they were generated at the same point as the cut. Moreover, if the call  $\text{r } 2 \text{ Z}$  itself produces multiple alternatives, only the first one is committed, while the remaining alternatives are discarded. This is because such alternatives have been created at a deeper depth in the search tree than the cut.

Concretely, the call  $\text{r } 2 \text{ Z}$  yields two solutions, assigning  $\text{Z}$  the values 4 and 8, respectively. The second solution is eliminated by the cut, and only the first assignment is preserved.

se metti  $r_1 = g A$   
 $B :- f A B$ . allora  
 $g e f$  sono funzioni, e puoi spiegare anche l'idea  
 del detcheck qui

## 3 Semantics intro

We propose two operational semantics for a logic program with cut. The two semantics are based on different syntaxes, the first syntax (called tree) exploits a tree-like structure and is ideal both to have a graphical view of its evolution while the state is being interpreted and to prove lemmas over it. The second syntax, called elpi, is the ELPI's syntax and has the advantage of reducing the computational cost of cutting and backtracking alternatives by using shared pointers. We aim to prove the equivalence of the two semantics together with some interesting lemmas of the cut behavior.

### 3.1 Tree semantics

```
Inductive tree :=
| KO | OK | TA : A -> tree
| Or : option tree -> Sigma -> tree -> tree
| And : tree -> seq A -> tree -> tree.
```

In the tree we distinguish 5 main cases: *KO*, *OK*, and are special meta-symbols representing, respectively, the failed and a successful terminal. These symbols are considered meta because they are internal intermediate symbols used to give structure to the tree.

The *TA* constructor (acronym for tree-atom) is the constructor of atoms in the tree.

TA = Todo/  
 Goal?

```

Fixpoint get_end s A : Sigma * tree:=
  match A with
  | TA _ | KO | OK => (s, A)
  | Or None s1 B => get_end s1 B
  | Or (Some A) _ _ => get_end s A
  | And A _ B =>
    let (s', pA) := get_end s A in
    if pA == OK then get_end s' B
    else (s', pA)
  end.

```

(a) Defintion of *get\_end*

104     The two recursive cases of a tree are the *Or* and *And* non-terminals. The *Or* non-terminal  
 105      $A \vee B_\sigma$  denotes a disjunction between two trees  $A$  and  $B$ . The first branch is optional, if  
 106     absent it represents a dead tree, i.e. a tree that has been entirely explored. The second  
 107     branch is annotated with a suspended substitution  $\sigma$  so that, upon backtracking to  $B$ ,  $\sigma$  is  
 108     used as the initial substitution for the execution of  $B$ .

109     The *And* non-terminal  $A \wedge_{B_0} B$  represents a conjunction of two trees  $A$  and  $B$ . We call  $B_0$   
 110     the reset point for  $B$ ; it is used to restore the state of  $B$  to its initial form if a backtracking  
 111     operation occurs on  $A$ . Intuitively, let  $t2l$  be the function flattening a tree in a list of sequents  
 112     disjunction, in PROLOG-like syntax the tree  $A \wedge_{B_0} B$  becomes  $(A_1, t2l B); (A_2, B_0); \dots; (A_n, B_0)$   
 t2l nope, mettiti  
 un r3 = 18  
 X Z :- r .., if  
 .., !, e rifatti  
 all'esempio dell'  
 sezione prima  
 (fai in modo che  
 f funzioni solo  
 con la seconda  
 regola per r)  
 associate to 122  
 the...  
 as much as 124  
 needed, indee  
 prolog programs  
 do not necessar 126  
 ily terminate 127  
 128  
 129  
 130  
 131  
 132  
 133  
 134

non-terminal è  
 roba di gram-  
 matiche, usa  
 nodes/con-  
 structors

A graphical representation of a tree is shown in Figure 4a. To make the graph more compact, the *And* and *Or* non-terminals are n-ary rather than binary, with right-binding priority. The *KO* terminal act as the neutral elements in the *Or* list, while *OK* is the neutral element of the *And* list.

The interpretation of a tree is performed by two main routines: *step* and *next\_alt* that traverse the tree depth-first, left-to-right. Then, then *run* inductive makes the transitive closure of step *step* and *next\_alt*: it iterates the calls to its auxiliary functions. In Types 7–9 we give the types contrats of these symbols where *fvS* is a set of variable names.

Inductive step\_tag := Expanded | CutBrothers | Failed | Success. (6)

Definition step : program -> fvS -> Sigma -> tree -> (fvS \* step\_tag \* tree) := (7)

Definition next\_alt : bool -> tree -> option tree := (8)

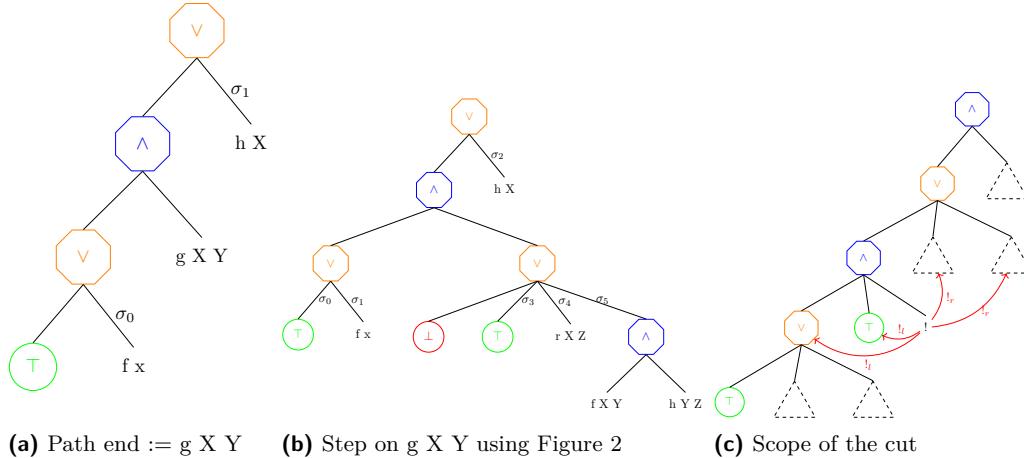
Inductive run (u:Unif) (p : program): fvS -> Sigma -> tree -> Sigma -> option tree -> Prop := (9)

The tree interpreter, as in prolog, explores the state in DFS strategy, to discover the substitution and the leaf of the tree that should be interpreted. The *get\_end* routine, shown in Figure 3a, accomplishes to this task. The *get\_end* returns its inputs if the tree is a leaf. Otherwise, if the tree is a disjunction, the path continues on the left subtree, if it exists, otherwise it recursively retrieves the wanted piece of information in the rhs using the substitution stored in the *Or* branch: the current substition when we cross the rhs of a *Or* is the one store in the *Or* node itself. In the case of a conjunction, if the to-be-explored leaf in the lhs is *OK*, then we look for the *get\_end* in the rhs, otherwise we return the result of the lhs.

We derive the following two functions from *get\_end*:

136     Definition get\_subst s A := (get\_end s A).1. (1)

137     Definition path\_end A := (get\_end empty A).2. (\*empty is the empty subst\*) (2)



**Figure 4** Some tree representations

138 In Figure 4a the *path\_end* of the tree is `g X Y`.

139 Below we define three special kinds of trees depending on their *path\_end*.

140 **Definition** `success A := path_end A == OK.` (3)

141 **Definition** `failed A := path_end A == KO.` (4)

142 **Definition** `path_atom A := if path_end A is TA _ then true else false.` (5)

143 The latter definition identifies path ending in an atom.

### 144 3.1.1 The *step* procedure

145 The *step* procedure takes as input a program, a set of free variables (*fv*), a substitution, and  
146 a tree, and returns an updated set of free variables, a *step\_tag*, and an updated tree.

147 Free variables are those variables that appear in a tree; they are used and updated when  
148 a backchaining operation takes place.

149 The *step\_tag* (see Type 6) indicates the kind of an internal tree step: `CutBrothers` denotes  
150 the interpretation of a superficial cut, i.e., a cut whose parent nodes are all *And*-nodes.  
151 `Expanded` denotes the interpretation of non-superficial cuts or predicate calls. `Failure` and  
152 `Success` are returned for, respectively, `failed` and `success` trees.

153 The step procedure is intended to interpretate atoms, that is, it transforms the tree iff its  
154 *path\_end* is an atom, otherwsise, it returns the identity.

155 **Lemma** `succ_step_iff u p fv s A: success A <-> step u p fv s A = (fv, Success, A)`. (1)  
156 **Lemma** `fail_step_iff u p fv s A: failed A <-> step u p fv s A = (fv, Failed, A)`. (2)

157 *Call step* The interpretation of a call *c* stars by calling the *bc* function on *c*. The output  
158 list *l* is taken to represent build the new subtree. If *l* is empty then *KO* tree is returned,  
159 otherwise the subtree is a right-skewed tree made of *n* inner *Or* nodes, where *n* is the length  
160 of *l*. The root has *KO* as left child. The lhs of the other nodes is a right-skewed tree of *And*  
161 nodes. The *And* nodes are again a right-seked tree containing premises of the selected rule .

162 A step in the tree of Figure 4a makes a backchain operation over the query `g X Y` and, in  
163 the program defined in Figure 2, the new tree would be the one in Figure 4b. We have put a  
164 red border aroung the new generated subtree. It is a disjunction of four subtrees: the first  
165 node is the *KO* node (by default), the second is *OK*, since *r1* has no premises, the third and  
166 the fourth contains the premises of respectively *r2* and *r3*.

dire dei reset  
point

dire che le sostituzioni del backchain sono importanti e

CVIT 2016

*Cut step* The cut case is delicate, since interpreting a cut in a tree has three main impacts: at first the cut is replaced by the *OK* node, then some special subtrees, in the scope of the *Cut*, are cut away: in particular we need to soft-kill the left-siblings of the *Cut* and hard-kill the right-uncles of the the *Cut*.

► **Definition 1** (Left-siblings (resp. right-sibling)). Given a node  $A$ , the left-siblings (resp. right-sibling) of  $A$  are the list of subtrees sharing the same parent of  $A$  and that appear on its left (resp. right).

► **Definition 2** (Right-uncles). Given a node  $A$ , the right-uncles of  $A$  are the list of right-sibling of the father of  $A$ .

► **Definition 3** (Hard-kill,  $!_r$ ). Given a tree  $t$ , hard-kill replaces the given subtree with the KO node

► **Definition 4** (Soft-kill,  $!\_l$ ). Given a successfull tree  $t$ , soft-kill replaces with KO all subtrees that are not part of the path in  $t$  leading to the OK node.

An example of the impact of the cut is shown in Figure 4c, the dashed triangles represent generic trees. The step routine interprets the cut since it is the node in its path-end: we pass through a and all trees on the left of the cut are successful. In the example we have 4 arrows tagged with the  $!_l$  or  $!_r$  symbols. The  $!_l$  arrows go left and soft-kill the pointed subtree, it keeps *OK* nodes since they are part of the tree leading to the cut, and replaces the other subtrees with *KO*. The  $!_r$  procedure replaces the nodes pointed by the arrows with *KO*.

### 186 3.1.2 The *next\_alt* procedure

187 It is evident that the *step* alone is not sufficient to reproduce entirely the behavior of the full  
 188 expected prolog interpreter. In particular, we need to bracktrack on failures. Moreover, in  
 189 case of success, we should return a state where the state is cleaned of the success itself, this  
 190 is essential to, non deterministically, find all the solution of a given query. By Lemmas 1  
 191 and 2, we know that *step* returns the identity on successful and failed states. In order to  
 192 continue the computation on these particular trees, we need the *next\_alt* procedure aiming  
 193 to expecially work with failed and successful trees: and its implementation in Figure 5.

The *next\_alt* procedure takes a boolean and a tree, clean it from failures or success and returns a new tree if this tree still contains a non explored path. The idea behind *next\_alt* is to clean recursively every subtree in DFS order if its *path\_end* is a failure. Moreover, if the boolean passed to *next\_alt* is true, then it erases the first successful path in the tree.

Some interesting property of *next\_alt* are shown below and allow to see how *next\_alt* complements *step*.

200 Lemma path\_atom\_next\_alt\_id b A: path\_atom A  $\rightarrow$  next\_alt b A = Some A. (3)

201 Lemma next\_alt\_failedF b A A': next\_alt b A = Some A'  $\rightarrow$  failed A' = false. (4)

For example, in Figure 4b the step procedure has created a failed state: its path-end ends in *KO*. The expected behavior of *next\_alt* is to take this *KO* node and make it a .... This allows *step* to continue the exploration of the tree. In particular, the path-end of this new state end in *OK*. The step leaves the state unchanged producing the new substitution. This solution however is not unique, we should be able to backtrack on this successful state. To do so we can call *next\_alt* and it will deadify the *OK* node allowing *step* to proceed on r X Z.

```

Definition next_alt : bool -> tree -> option tree :=
fix next_alt b A :=
  match A with
  | KO => None
  | OK => if b then None else Some OK
  | TA _ => Some A
  | And A BO B =>
    let build_BO A := And A BO (big_and BO) in
    if success A then
      match next_alt b B with
      | None => omap build_BO (next_alt true A)
      | Some B' => Some (And A BO B')
    end
    else if failed A then omap build_BO (next_alt false A)
    else Some (And A BO B)
  | Or None sB B => omap (fun x => Or None sB x) (next_alt b B)
  | Or (Some A) sB B =>
    match next_alt b A with
    | None => omap (fun x => Or None sB x) (next_alt false B)
    | Some A' => Some (Or (Some A') sB B)
  end
end.

```

■ **Figure 5** *next\_alt* implementation

### 208 3.1.3 The *run* inductive

209 The run inductive is nothing but the transitive closure of the *step* and *next\_alt*

$$\begin{array}{c}
 \frac{\text{success} A}{\text{run } fv s_1 A (\text{get\_subst } s_1 A) (\text{next\_alt } \perp A)} \text{ run\_done} \\
 \frac{\text{path\_atom } A \quad \text{step } u p fv_0 s_1 A = (fv_1, \text{ st, } B) \quad \text{run } fv_1 s_1 B s_2 r}{\text{run } fv_0 s_1 A s_2 r} \text{ run\_step} \\
 \frac{\text{failed } A \quad \text{next\_alt } \perp A = \text{Some } B \quad \text{run } fv_0 s_1 B s_2 r}{\text{run } fv_0 s_1 A s_2 r} \text{ run\_fail}
 \end{array}$$

### 213 3.1.4 Valid tree

214 Reasoning on a the tree semantics allows to identify an invariant that

## 215 3.2 Elpi semantics

216 TODO: dire che la semantica ad albero è più facile per le prove

217 The ELPI interpreter is based on an operational semantics close to the one picked by  
 218 Pusch in [16], in turn closely related to the one given by Debray and Mishra in [6, Section  
 219 4.3]. Push mechanized the semantics in Isabelle/HOL together with some optimizations that  
 220 are present in the Warren Abstract Machine [20, 1].

221 In these operational semantics we need to decorate the cut atom with a list of alternative,  
 222 morally a pointer to a sub-list of the overall alternatives. An atom in the elpi semantics is  
 223 defined as follows:

```
Inductive alts := no_alt | more_alt of (Sigma * goals) & alts
with goals := no_goals | more_goals of (A * alts) & goals .
```

We are completely loosing the tree structure. There are no clean reset points. The backtracking operation is simpler: it is the tail function. The cutr and cutl operations disappears: the alternatives are stored directly in the cutE terminal.

227 The elpi interpreter is as follows:

```
(*TODO: add system of rules*)
Inductive nur : Sigma -> goals -> alts -> Sigma -> alts -> Type :=
| StopE s a : nur s nilC a s a
| CutE s s1 a ca r gl : nur s gl ca s1 r -> nur s ((cutE ca) :: gl) a s1 r
| CallE p s s1 a b bs gl r t :
  F u p t s = [:: b & bs ] ->
    nur b.1 (save_goals a gl (a2gs1 p b)) (save_alts a gl ((aa2gs p) bs) ++ a) s1 r ->
    nur s ((callE p t) :: gl) a s1 r
| FailE p s s1 s2 t gl a al r :
  F u p t s = [::] -> nur s1 a al s2 r -> nur s ((callE p t) :: gl) ((s1, a) :: al) s2 r.
```

228 The translation of a tree to a list is as follows:

```

Fixpoint t21 (A: tree) s (bt : alts) : alts :=

  match A with
  | OK           => [:: (s, [::]) ]
  | KO           => [::]
  | TA a         => [:: (s, [:: (a,[::]) ])]
  | Or None s1 B => add_ca_deep bt (t21 B s1 [::])
  | Or (Some A) s1 B   =>
    let lB := t21 B s1 [::] in
    let lA := t21 A s lB in
    add_ca_deep bt (lA ++ lB)
  | And A B0 B   =>
    let lB0 := r21 B0 in
    let lA := t21 A s bt in
    if lA is [:: (s1A, x) & xs] then
      let xz := add_deepG bt lB0 x in
      let xs := add_deep bt lB0 xs in
      let xs := map (catr lB0) xs in
      let lB := t21 B s1A (xs ++ bt) in
      (map (catl xz) lB) ++ xs
    else [::]
  end.

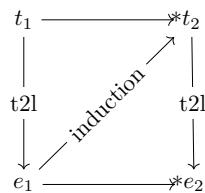
```

### ► Theorem 5 (tree to elpi).

```

229           $\forall A \sigma_1 B \sigma_2 b \sigma_0, \text{vt } A \rightarrow$ 
230           $\text{run}_u \sigma_1 A (\text{Some } \sigma_2) B b \rightarrow$ 
231           $\exists x xs, \text{t2l } A \sigma_1 \emptyset = x ::: xs \wedge \text{nur}_u x.1 x.2 xs \sigma_2 (\text{t2l } B \sigma_0 \emptyset).$ 

```



**Figure 6** Induction scheme for Theorem 6

► **Theorem 6** (`elpi_to_tree`).

```

232                                      $\forall \sigma_1 \sigma_2 a \; na \; g,$ 
233                                      $\text{nur}_u \sigma_1 g \; a \; \sigma_2 \; na \rightarrow$ 
234                                      $\forall \sigma_0 t, \text{vt} \; t \rightarrow (\text{t2l} \; t \; \sigma_0 \emptyset) = ((\sigma_1, g) :: a) \rightarrow$ 
235                                      $\exists t' \; n, \text{run}_u \sigma_0 \; t \; (\text{Some} \; \sigma_2) \; t' \; n \wedge \text{t2l} \; t' \; \sigma_0 \emptyset = na.$ 

```

The proof of Theorem 6 is based on the idea explained in [2, Section 3.3]. An ideal statement for this lemma would be: given a function  $12t$  transforming an elpi state to a tree, we would have that the execution of an elpi state  $e$  is the same as executing  $run$  on the tree resulting from  $12t(e)$ . However, it is difficult to retrieve the structure of an elpi state and create a tree from it. This is because, in an elpi state, we have no clear information about the scope of an atom inside the list and, therefore, no evident clue about where this atom should be placed in the tree.

Our theorem states that, starting from a valid state  $t$  which translates to a list of alternatives  $(\sigma_1, g) :: a$ . If we run in elpi the list of alternatives, then the execution of the tree  $t$  returns the same result as the execution in elpi. The proof is performed by induction on the derivations of the elpi execution. We have 4 derivations.

247 We have 4 case to analyse:

248 — References

- 249 1 Hassan Aït-Kaci. *Warren’s Abstract Machine: A Tutorial Reconstruction*. The MIT Press, 08  
 250 1991. doi:[10.7551/mitpress/7160.001.0001](https://doi.org/10.7551/mitpress/7160.001.0001).

251 2 Yves Bertot. A certified compiler for an imperative language. Technical Report RR-3488,  
 252 INRIA, September 1998. URL: <https://hal.inria.fr/inria-00073199v1>.

253 3 Valentin Blot, Denis Cousineau, Enzo Crance, Louise Dubois de Prisque, Chantal Keller,  
 254 Assia Mahboubi, and Pierre Vial. Compositional pre-processing for automated reasoning in  
 255 dependent type theory. In Robbert Krebbers, Dmitriy Traytel, Brigitte Pientka, and Steve  
 256 Zdancewic, editors, *Proceedings of the 12th ACM SIGPLAN International Conference on  
 257 Certified Programs and Proofs, CPP 2023, Boston, MA, USA, January 16–17, 2023*, pages  
 258 63–77. ACM, 2023. doi:[10.1145/3573105.3575676](https://doi.org/10.1145/3573105.3575676).

259 4 Cyril Cohen, Enzo Crance, and Assia Mahboubi. Trocq: Proof transfer for free, with or  
 260 without univalence. In Stephanie Weirich, editor, *Programming Languages and Systems*, pages  
 261 239–268, Cham, 2024. Springer Nature Switzerland.

262 5 Cyril Cohen, Kazuhiko Sakaguchi, and Enrico Tassi. Hierarchy Builder: Algebraic hierarchies  
 263 Made Easy in Coq with Elpi. In *Proceedings of FSCD*, volume 167 of *LIPICS*, pages 34:1–34:21,  
 264 2020. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.FSCD.2020.34>.  
 265 doi:[10.4230/LIPIcs.FSCD.2020.34](https://doi.org/10.4230/LIPIcs.FSCD.2020.34).

266 6 Saumya K. Debray and Prateek Mishra. Denotational and operational semantics for prolog. *J.  
 267 Log. Program.*, 5(1):61–91, March 1988. doi:[10.1016/0743-1066\(88\)90007-6](https://doi.org/10.1016/0743-1066(88)90007-6).

## 23:10 Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

- 268 7 Cvetan Dunchev, Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. ELPI: fast,  
269 embeddable,  $\lambda$ Prolog interpreter. In *Proceedings of LPAR*, volume 9450 of *LNCS*, pages  
270 460–468. Springer, 2015. URL: <https://inria.hal.science/hal-01176856v1>, doi:10.1007/978-3-662-48899-7\\_32.
- 272 8 Davide Fissore and Enrico Tassi. A new Type-Class solver for Coq in Elpi. In *The Coq  
273 Workshop*, July 2023. URL: <https://inria.hal.science/hal-04467855>.
- 274 9 Davide Fissore and Enrico Tassi. Higher-order unification for free!: Reusing the meta-  
275 language unification for the object language. In *Proceedings of PPDP*, pages 1–13. ACM, 2024.  
276 doi:10.1145/3678232.3678233.
- 277 10 Davide Fissore and Enrico Tassi. Determinacy checking for elpi: an higher-order logic program-  
278 ming language with cut. In *Practical Aspects of Declarative Languages: 28th International  
279 Symposium, PADL 2026, Rennes, France, January 12–13, 2026, Proceedings*, pages 77–95,  
280 Berlin, Heidelberg, 2026. Springer-Verlag. doi:10.1007/978-3-032-15981-6\_5.
- 281 11 Benjamin Grégoire, Jean-Christophe Léchenet, and Enrico Tassi. Practical and sound equality  
282 tests, automatically. In *Proceedings of CPP*, page 167–181. Association for Computing  
283 Machinery, 2023. doi:10.1145/3573105.3575683.
- 284 12 Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. Implementing type theory  
285 in higher order constraint logic programming. In *Mathematical Structures in Computer  
286 Science*, volume 29, pages 1125–1150. Cambridge University Press, 2019. doi:10.1017/  
287 S0960129518000427.
- 288 13 Robbert Krebbers, Luko van der Maas, and Enrico Tassi. Inductive Predicates via Least  
289 Fixpoints in Higher-Order Separation Logic. In Yannick Forster and Chantal Keller, editors,  
290 *16th International Conference on Interactive Theorem Proving (ITP 2025)*, volume 352 of *Leib-  
291 niz International Proceedings in Informatics (LIPIcs)*, pages 27:1–27:21, Dagstuhl, Germany,  
292 2025. Schloss Dagstuhl – Leibniz-Zentrum für Informatik. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs. ITP.2025.27>, doi:10.4230/LIPIcs. ITP.2025.27.
- 294 14 Dale Miller. A logic programming language with lambda-abstraction, function variables, and  
295 simple unification. In *Extensions of Logic Programming*, pages 253–281. Springer, 1991.
- 296 15 Dale Miller and Gopalan Nadathur. *Programming with Higher-Order Logic*. Cambridge  
297 University Press, 2012.
- 298 16 Cornelia Pusch. Verification of compiler correctness for the wam. In Gerhard Goos, Juris  
299 Hartmanis, Jan van Leeuwen, Joakim von Wright, Jim Grundy, and John Harrison, editors,  
300 *Theorem Proving in Higher Order Logics*, pages 347–361, Berlin, Heidelberg, 1996. Springer  
301 Berlin Heidelberg.
- 302 17 Enrico Tassi. Elpi: an extension language for Coq (Metaprogramming Coq in the Elpi  $\lambda$ Prolog  
303 dialect). In *The Fourth International Workshop on Coq for Programming Languages*, January  
304 2018. URL: <https://inria.hal.science/hal-01637063>.
- 305 18 Enrico Tassi. Deriving proved equality tests in Coq-Elpi. In *Proceedings of ITP*, volume 141 of  
306 *LIPIcs*, pages 29:1–29:18, September 2019. URL: <https://inria.hal.science/hal-01897468>,  
307 doi:10.4230/LIPIcs.CVIT.2016.23.
- 308 19 Luko van der Maas. Extending the Iris Proof Mode with inductive predicates using Elpi.  
309 Master’s thesis, Radboud University Nijmegen, 2024. doi:10.5281/zenodo.12568604.
- 310 20 David H.D. Warren. An Abstract Prolog Instruction Set. Technical Report Technical Note 309,  
311 SRI International, Artificial Intelligence Center, Computer Science and Technology Division,  
312 Menlo Park, CA, USA, October 1983. URL: <https://www.sri.com/wp-content/uploads/2021/12/641.pdf>.