Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

Jane Open Access ✉ 🏠 🆔

Dummy University Computing Laboratory, [optional: Address], Country

My second affiliation, Country

Joan R. Public[1] ✉ 🆔

Department of Informatics, Dummy College, [optional: Address], Country

──── **Abstract** ────────────────────────────

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Praesent convallis orci arcu, eu mollis dolor. Aliquam eleifend suscipit lacinia. Maecenas quam mi, porta ut lacinia sed, convallis ac dui. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Suspendisse potenti.

1 Introduction

ELPI is a dialect of λPROLOG (see [14, 15, 7, 12]) used as an extension language for the ROCQ prover (formerly the COQ proof assistant). ELPI has become an important infrastructure component: several projects and libraries depend on it [13, 3, 4, 19, 8, 9]. Examples include the Hierarchy-Builder library-structuring tool [5] and Derive [17, 18, 11], a program-and-proof synthesis framework with industrial applications at SkyLabs AI.

Starting with version 3, ELPI gained a static analysis for determinacy [10] to help users tame backtracking. ROCQ users are familiar with functional programming but not necessarily with logic programming and uncontrolled backtracking is a common source of inefficiency and makes debugging harder. The determinacy checkers identifies predicates that behave like functions, i.e., predicates that commit to their first solution and leave no *choice points* (places where backtracking could resume).

This paper reports our first steps towards a mechanization, in the ROCQ prover, of the determinacy analysis from [10]. We focus on the control operator *cut*, which is useful to restrict backtracking but makes the semantic depart from a pure logical reading.

We formalize two operational semantics for PROLOG with cut. The first is a stack-based semantics that closely models ELPI's implementation and is similar to the semantics mechanized by Pusch in ISABELLE/HOL [16] and to the model of Debray and Mishra [6, Sec. 4.3]. This stack-based semantics is a good starting point to study further optimizations used by standard PROLOG abstract machines [20, 1], but it makes reasoning about the scope of *cut* difficult. To address that limitation we introduce a tree-based semantics in which the branches pruned by *cut* are explicit and we prove the two semantics equivalent. Using the

---

[1] Optional footnote, e.g. to mark corresponding author

```
Inductive P := IP of nat. Inductive D := ID of nat. Inductive V := IV of nat.

Inductive Tm :=                          Inductive Callable :=
  | Tm_P of P     | Tm_D   of D            | Callable_P   of P
  | Tm_V of V     | Tm_App of Tm & Tm.     | Callable_App of Callable & Tm.
```

🟨 **Figure 1** Tm and Callable types

40   tree-based semantics we then show that if every rule of a predicate passes the determinacy
41   analysis, the call to a deterministic predicate does not leave any choice points.

2   Common code: the language

put unif and pro-
gram in variable,
hides from types

43   Before going to the two semantcis, we show the piece of data structure that are shared by
44   the them. The smallest unit of code that we can use in the langauge is an atom. The atom
45   inductive (see Type 1) is either a cut or a call. A call carries a callable term (see Figure 1).
46   A term (Tm) is either a predicate, a datum, a variable or the binary application of a term to
47   another. A Callable is a term accepting predicates only predicates as functors.

48   `Inductive A := cut | call : Callable -> A.`   (1)
49   `Record R := mkR { head : Callable; premises : list A }.`   (2)
50   `Record program := { rules : seq R; sig : sigT }.`   (3)
51   `Definition Sigma := {fmap V -> Tm}.`   (4)
52   `Definition bc : Unif -> program -> fvS -> Callable ->`
     `                  Sigma -> fvS * seq (Sigma * R) :=`   (5)

53   A rule (see Type 2) is made a head of type term and a list of premises, the premises are
54   atoms. A program (see Type 3) is made by a list of rules and a mapping from predicates to
55   their signatures. The type sigT is the classic type from the simply typed lambda calculus, i.e.
56   it is either a base type or an arrow. We decorate arrows to know the mode of the lhs type.
57   A substitution (see Type 4) is a mapping from variables to terms. It is the output of a
58   successful query and is often called the output of a query.

```
Record Unif := {
  unify : Tm -> Tm -> Sigma -> option Sigma;
  matching : Tm -> Tm -> Sigma -> option Sigma;
}.
```

59   The backchain function (bc, see Type 5) filters the rules in the program that can be
60   used on a given query. It takes: a unifcator U which explains how to unify terms up to
61   standard unifcation (for output terms) or matching (for input terms); a program P to explore
62   and filter; a set S of free variable (fvS) allowing to fresh the program P by renaming the
63   its variables; a query q; and the substitution σ in which the query q lives. The result of a
64   backchain operation is couple made of an extension of S containing the new variables that
65   have been allocated during the unification phase and a list of filtered rules r accompagnate
66   by their a subistution. This substitution is the result of the unification of q with the head of
67   each rule in r.
68   In Figure 2, we have an example of a simple ELPI program which will be used in the
69   following section of the paper as an example to show how backtracking and the cut operator
70   works in the semantcis we propose. The translation of these rules in the ROCQ representation
71   is straightfoword.

```
f 1 2.   f 2 3.   r 2 4.   r 2 8.
g X X.                      % r1
g X Z :- r X Z, !.          % r2
g X Z :- f X Y, f Y Z.      % r3
```

🟨 **Figure 2** Small ELPI program example

2.1   The cut operator

The semantics of the cut operator adopted in the ELPI language corresponds to the *hard cut* operator of standard SWI-PROLOG. This operator has two primary purposes. First, it eliminates all alternatives that are created either simultaneously with, or after, the introduction of the cut into the execution state.

To illustrate this high-level description, consider the program shown in Figure 2 and the query $q = $ `g 2 Z`. All three rules for `g` can be used on the query q. They are tried according to their order of appearance in the program: rule r_1 is tried first, followed by r_2, and r_3.

The first rule has no premises and immediately returns the assignment `Z = 2`. However, the computation does not terminate at this point, since two additional unexplored alternatives remain, corresponding to the premises of rules r_2 and r_3.

The premises of rule r_2 are `r 2 Z, !`. At this stage, the role of the cut becomes apparent. If the premise `r 2 Z` succeeds, the cut commits to this choice and removes the premises of rule r_3 from the alternative list, as they were generated at the same point as the cut. Moreover, if the call `r 2 Z` itself produces multiple alternatives, only the first one is committed, while the remaining alternatives are discarded. This is because such alternatives have been created at a deeper depth in the search tree than the cut.

Concretely, the call `r 2 Z` yields two solutions, assigning `Z` the values 4 and 8, respectively. The second solution is eliminated by the cut, and only the first assignment is preserved.

3   Semantics intro

We propose two operational semantics for a logic program with cut. The two semantics are based on different syntaxes, the first syntax (called tree) exploits a tree-like structure and is ideal both to have a graphical view of its evolution while the state is being intepreted and to prove lemmas over it. The second syntax, called elpi, is the ELPI's syntax and has the advantage of reducing the computational cost of cutting and backtracking alternatives by using shared pointers. We aim to prove the equivalence of the two semantics together with some interesting lemmas of the cut behavior.

3.1   Tree semantics

```
Inductive tree :=
  | KO | OK | TA : A -> tree
  | Or  : option tree -> Sigma -> tree -> tree
  | And : tree -> seq A -> tree -> tree.
```

In the tree we distinguish 5 main cases: *KO*, *OK*, and are special meta-symbols representing, respectively, the failed and a successful terminal. These symbols are considered meta because they are internal intermediate symbols used to give structure to the tree.

The *TA* constructor (acronym for tree-atom) is the constructor of atoms in the tree.

```
Fixpoint get_end s A : Sigma * tree:=
  match A with
  | TA _ | KO | OK => (s, A)
  | Or None s1 B => get_end s1 B
  | Or (Some A) _ _ => get_end s A
  | And A _ B =>
    let (s', pA) := get_end s A in
    if pA == OK then get_end s' B
    else (s', pA)
  end.
```

(a) Defintion of *get_end*

$_{104}$ The two recursive cases of a tree are the *Or* and *And* non-terminals. The *Or* non-terminal
$_{105}$ $A \vee B_\sigma$ denotes a disjunction between two trees A and B. The first branch is optional, if
$_{106}$ absent it represents a dead tree, i.e. a tree that has been entirely explored. The second
$_{107}$ branch is annotated with a suspended substitution σ so that, upon backtracking to B, σ is
$_{108}$ used as the initial substitution for the execution of B.

$_{109}$ The *And* non-terminal $A \wedge_{B_0} B$ represents a conjunction of two trees A and B. We call B_0
$_{110}$ the reset point for B; it is used to restore the state of B to its initial form if a backtracking
$_{111}$ operation occurs on A. Intuitively, let *t2l* be the function flattening a tree in a list of sequents
$_{112}$ disjnction, in PROLOG-like syntax the tree $A \wedge_{B_0} B$ becomes $(A_1, t2l\ B); (A_2, B_0); \ldots; (A_n, B_0)$
$_{113}$ where $t2l(A) = A_1, \ldots, A_n$.

$_{114}$ A graphical representation of a tree is shown in Figure 4a. To make the graph more
$_{115}$ compact, the *And* and *Or* non-terminals are n-ary rather than binary, with right-binding
$_{116}$ priority. The *KO* terminal act as the neutral elements in the *Or* list, while *OK* is the neutral
$_{117}$ element of the *And* list.

$_{118}$ The interpretation of a tree is performed by two main routines: *step* and *next_alt* that
$_{119}$ traverse the tree depth-first, left-to-right. Then, then *run* inductive makes the transitive
$_{120}$ closure of step *step* and *next_alt*: it iterates the calls to its auxiliary functions. In Types 7–9
$_{121}$ we give the types contrats of these symbols where `fvS` is a set of variable names.

$_{122}$ `Inductive step_tag := Expanded | CutBrothers | Failed | Success.`   (6)

$_{123}$ `Definition step : program -> fvS -> Sigma -> tree -> (fvS * step_tag * tree) :=` (7)

$_{124}$ `Definition next_alt : bool -> tree -> option tree :=`   (8)
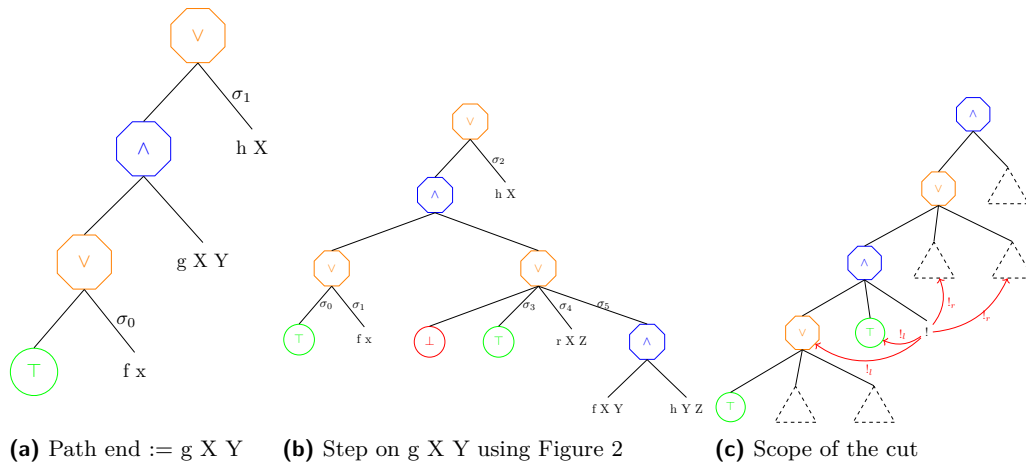
$_{125}$ `Inductive run (u:Unif) (p : program): fvS -> Sigma -> tree ->`
            `Sigma -> option tree -> Prop :=`   (9)

$_{126}$ The tree interpreter, as in prolog, explores the state in DFS strategy, to discover the
$_{127}$ substitution and the leaf of the tree that should be interpreted. The *get_end* routine, shown
$_{128}$ in Figure 3a, accomplishes to this task. The *get_end* returns its inputs if the tree is a
$_{129}$ leaf. Otherwise, if the tree is a disjunction, the path continues on the left subtree, if it
$_{130}$ exists, otherwise it recursively retrieves the wanted piece of information in the rhs using the
$_{131}$ substitution stored in the *Or* branch: the current substition when we cross the rhs of a *Or* is
$_{132}$ the one store in the Or node itself. In the case of a conjunction, if the to-be-explored leaf in
$_{133}$ the lhs is *OK*, then we look for the *get_end* in the rhs, otherwise we return the result of the
$_{134}$ lhs.

$_{135}$ We derive the following two functions from *get_end*:

$_{136}$ `Definition get_subst s A := (get_end s A).1.`   (1)

$_{137}$ `Definition path_end A := (get_end empty A).2.` *(* empty is the empty subst *)*   (2)

(a) Path end := g X Y **(b)** Step on g X Y using Figure 2 **(c)** Scope of the cut

Figure 4 Some tree representations

In Figure 4a the *path_end* of the tree is `g X Y`.

Below we define three special kinds of trees depending on their *path_end*.

`Definition success A := path_end A == OK.` (3)

`Definition failed A := path_end A == KO.` (4)

`Definition path_atom A := if path_end A is TA _ then true else false.` (5)

The latter definition identifies path ending in an atom.

3.1.1 The *step* procedure

The *step* procedure takes as input a program, a set of free variables (fv), a substitution, and a tree, and returns an updated set of free variables, a *step_tag*, and an updated tree.

Free variables are those variables that appear in a tree; they are used and updated when a backchaining operation takes place.

The *step_tag* (see Type 6) indicates the kind of an internal tree step: `CutBrothers` denotes the interpretation of a superficial cut, i.e., a cut whose parent nodes are all *And*-nodes. `Expanded` denotes the interpretation of non-superficial cuts or predicate calls. `Failure` and `Success` are returned for, respectively, *failed* and *success* trees.

The step procedure is intended to interpretate atoms, that is, it transforms the tree iff its *path_end* is an atom, otherwsise, it returnes the identity.

`Lemma succ_step_iff u p fv s A: success A <-> step u p fv s A = (fv, Success, A)` (1)

`Lemma fail_step_iff u p fv s A: failed A <-> step u p fv s A = (fv, Failed, A).` (2)

Call step The interpretation of a call *c* stars by calling the *bc* function on *c*. The output list *l* is taken to represent build the new subtree. If *l* is empty then *KO* tree is returned, otherwise the subtree is a right-skewed tree made of *n* inner *Or* nodes, where *n* is the length of *l*. The root has *KO* as left child. The lhs of the other nodes is a right-skewed tree of *And* nodes. The *And* nodes are again a right-seked tree containing premises of the selected rule .

A step in the tree of Figure 4a makes a backchain operation over the query `g X Y` and, in the program defined in Figure 2, the new tree would be the one in Figure 4b. We have put a red border aroung the new generated subtree. It is a disjunction of four subtrees: the first node is the *KO* node (by default), the second is *OK*, since r1 has no premises, the third and the fourth contains the premises of respectively r_2 and r_3.

dire dei reset point

dire che le sostituzioni del backchain sono importanti e

167   *Cut step* The cut case is delicate, since interpreting a cut in a tree has three main impacts:
168   at first the cut is replaced by the *OK* node, then some special subtrees, in the scope of the
169   *Cut*, are cut away: in particular we need to soft-kill the left-siblings of the *Cut* and hard-kill
170   the right-uncles of the the *Cut*.

171   ▶ **Definition 1** (Left-siblings (resp. right-sibling)). *Given a node A, the left-siblings (resp.*
172   *right-sibling) of A are the list of subtrees sharing the same parent of A and that appear on*
173   *its left (resp. right).*

174   ▶ **Definition 2** (Right-uncles). *Given a node A, the right-uncles of A are the list of right-sibling*
175   *of the father of A.*

176   ▶ **Definition 3** (Hard-kill, $!_r$). *Given a tree t, hard-kill replaces the given subtree with the*
177   KO *node*

178   ▶ **Definition 4** (Soft-kill, $!_l$). *Given a successfull tree t, soft-kill replaces with* KO *all subtrees*
179   *that are not part of the path in t leading to the* OK *node.*

180   An example of the impact of the cut is show in Figure 4c, the dashed triangles represent
181   generic trees. The step routine interprets the cut since it is the node in its path-end: we pass
182   throgh a and and all trees on the left of the cut are successful. In the example we have 4
183   arrow tagged with the $!_l$ or $!_r$ symbols. The $!_l$ arrows go left and soft-kill the pointed subtree,
184   it keeps *OK* nodes since they are part of the tree leading to the cut, and replaces the other
185   subtrees with *KO*. The $!_r$ procedure replaces the nodes pointed by the arrows with *KO*.

186   3.1.2   The *next_alt* procedure

187   It is evident that the *step* alone is not sufficient to reproduce entirely the behavior of the
188   full ELPI solver. In particular, *step* does not perform any backtracking nor it cleans success.
189   From Lemmas 1 and 2, *step* returns the identity. In order to continue the computation on
190   these particular trees, we need the *next_alt* procedure: its signature is provided in Type 8
191   and its implementation in Figure 5.

192   The *next_alt* procedure takes a boolean and a tree and return a new tree if it still contains
193   an alternative. The idea behind *next_alt* is to clean recursively remove every subtree in DFS
194   order if its *path_end* is a failure. Moreover, if the boolean passed to *next_alt* is true, then it
195   erases the first successful path in the tree.

196   An interesting property of *next_alt* is the following, allowing to see how *next_alt* comple-
197   ments *step*.

198   `Lemma path_atom_next_alt_id b A: path_atom A -> next_alt b A = Some A.`   (3)

199   For example, in Figure 4b the step procedure has created a failed state: its path-end ends
200   in *KO*. The expected behavior of *next_alt* is to take this *KO* node and make it a .... This
201   allows *step* to continue the exploration of the tree. In particular, the path-end of this new
subst taken form 202   state end in *OK*. The step leaves the state unchanged producing the new substitution. This
the or   203   solution however is not unique, we should be able to backtrack on this successful state. To do
204   so we can call *next_alt* and it will deadify the *OK* node allowing *step* to proceed on r X Z.

205   3.1.3   The *run* inductive

206   The run inductive is nothing but the transitive closure of the *step* and *next_alt*

$$207 \qquad \frac{success A}{run\ fv\ s_1\ A\ (get\_subst\ s_1\ A)\ (next\_alt\ \top\ A)}\ run\_done$$

```
Definition next_alt : bool -> tree -> option tree :=
  fix next_alt b A :=
  match A with
  | KO => None
  | OK => if b then None else Some OK
  | TA _ => Some A
  | And A B0 B =>
    let build_B0 A := And A B0 (big_and B0) in
    if success A then
      match next_alt b B with
      | None => omap build_B0 (next_alt true A)
      | Some B => Some (And A B0 B)
      end
    else if failed A then omap build_B0 (next_alt false A)
    else Some (And A B0 B)
  | Or None sB B => omap (fun x => Or None sB x) (next_alt b B)
  | Or (Some A) sB B =>
      match next_alt b A with
      | None => omap (fun x => Or None sB x) (next_alt false B)
      | Some nA => Some (Or (Some nA) sB B)
      end
  end.
```

■ Figure 5 *next_alt* implementation

3.1.4 Valid tree

Reasoning on a the tree semantics allows to identify an invariant that

3.2 Elpi semantics

TODO: dire che la semantica ad albero è puù faicle per le prove

The ELPI interpreter is based on an operational semantics close to the one picked by Pusch in [16], in turn closely related to the one given by Debray and Mishra in [6, Section 4.3]. Push mechanized the semantics in Isabelle/HOL together with some optimizations that are present in the Warren Abstract Machine [20, 1].

In these operational semantics we need to decorate the cut atom with a list of alternative, morally a pointer to a sub-list of the overall alternatives. An atom in the elpi semantcis is defined as follows:

```
Inductive alts :=
  | no_alt
  | more_alt : (Sigma * goals) -> alts -> alts
with goals :=
  | no_goals
  | more_goals : (A * alts) -> goals -> goals .
```

We are completely loosing the tree structure. There are no clean reset points. The backtracking operation is simpler: it is the tail function. The cutr and cutl operations disappears: the alternatives are stored directly in the cutE terminal.

The elpi interpreter is as follows:

```
(*TODO: add system of rules*)
Inductive nur : Sigma -> goals ->  alts -> Sigma -> alts -> Type :=
  | StopE s a : nur s nilC a s a
  | CutE s s1 a ca r gl : nur s gl ca s1 r -> nur s ((cutE ca) ::: gl) a s1 r
  | CallE p s s1 a b bs gl r t :
    F u p t s = [:: b & bs ] ->
      nur b.1 (save_goals a gl (a2gs1 p b)) (save_alts a gl ((aa2gs p) bs) ++ a) s1 r ->
        nur s ((callE p t) ::: gl) a s1 r
  | FailE p s s1 s2 t gl a al r :
    F u p t s = [::] -> nur s1 a al s2 r -> nur s ((callE p t) ::: gl) ((s1, a) ::: al) s2 r.
```

223   The translation of a tree to a list is as follows:

```
Fixpoint t2l (A: tree) s (bt : alts) : alts :=
  match A with
  | OK           => [:: (s, [::]) ]
  | KO           => [::]
  | TA a         => [:: (s, [:: (a,[::]) ]) ]
  | Or None s1 B => add_ca_deep bt (t2l B s1 [::])
  | Or (Some A) s1 B    =>
      let lB := t2l B s1 [::] in
      let lA := t2l A s lB in
      add_ca_deep bt (lA ++ lB)
  | And A B0 B    =>
      let lB0 := r2l B0 in
      let lA   := t2l A s bt in
      if lA is [:: (slA, x) & xs] then
        let xz := add_deepG bt lB0 x in
        let xs := add_deep bt lB0 xs in
        let xs := map (catr lB0) xs in
        let lB := t2l B slA (xs ++ bt) in
        (map (catl xz) lB) ++ xs
      else [::]
  end.
```
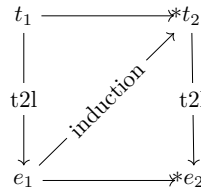
▶ **Theorem 5** (tree_to_elpi).

224
$$\forall A \ \sigma_1 \ B \ \sigma_2 \ b \ \sigma_0, \text{vt } A \to$$

225
$$\text{run}_u \ \sigma_1 \ A \ (Some \ \sigma_2) \ B \ b \to$$

226
$$\exists x \ xs, \text{t2l } A \ \sigma_1 \ \varnothing = x ::: xs \land \text{nur}_u \ x.1 \ x.2 \ xs \ \sigma_2 \ (\text{t2l } B \ \sigma_0 \ \varnothing).$$

▶ **Theorem 6** (elpi_to_tree).

227
$$\forall \sigma_1 \ \sigma_2 \ a \ na \ g,$$

228
$$\text{nur}_u \ \sigma_1 \ g \ a \ \sigma_2 \ na \to$$

229
$$\forall \sigma_0 \ t, \text{vt } t \to (\text{t2l } t \ \sigma_0 \ \varnothing) = ((\sigma_1, g) ::: a) \to$$

230
$$\exists t' \ n, \text{run}_u \ \sigma_0 \ t \ (Some \ \sigma_2) \ t' \ n \land \text{t2l } t' \ \sigma_0 \ \varnothing = na.$$

231   The proof of Theorem 6 is based on the idea explained in [2, Section 3.3]. An ideal
232 statement for this lemma would be: given a function `l2t` transforming an elpi state to a tree,
233 we would have have that the the execution of an elpi state e is the same as executing *run* on
234 the tree resulting from `l2t(e)`. However, it is difficult to retrive the strucutre of an elpi state

Figure 6 Induction scheme for Theorem 6

and create a tree from it. This is because, in an elpi state, we have no clear information about the scope of an atom inside the list and, therefore, no evident clue about where this atom should be place in the tree.

Our theorem states that, starting from a valid state t which translates to a list of alternatives $(\sigma_1, g) :: a$. If we run in elpi the list of alternatives, then the execution of the tree t returns the same result as the execution in elpi. The proof is performed by induction on the derivations of the elpi execution. We have 4 derivations.

We have 4 case to analyse:

References

1 Hassan Aït-Kaci. *Warren's Abstract Machine: A Tutorial Reconstruction.* The MIT Press, 08 1991. `doi:10.7551/mitpress/7160.001.0001`.

2 Yves Bertot. A certified compiler for an imperative language. Technical Report RR-3488, INRIA, September 1998. URL: `https://inria.hal.science/inria-00073199v1`.

3 Valentin Blot, Denis Cousineau, Enzo Crance, Louise Dubois de Prisque, Chantal Keller, Assia Mahboubi, and Pierre Vial. Compositional pre-processing for automated reasoning in dependent type theory. In Robbert Krebbers, Dmitriy Traytel, Brigitte Pientka, and Steve Zdancewic, editors, *Proceedings of the 12th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2023, Boston, MA, USA, January 16-17, 2023*, pages 63–77. ACM, 2023. `doi:10.1145/3573105.3575676`.

4 Cyril Cohen, Enzo Crance, and Assia Mahboubi. Trocq: Proof transfer for free, with or without univalence. In Stephanie Weirich, editor, *Programming Languages and Systems*, pages 239–268, Cham, 2024. Springer Nature Switzerland.

5 Cyril Cohen, Kazuhiko Sakaguchi, and Enrico Tassi. Hierarchy Builder: Algebraic hierarchies Made Easy in Coq with Elpi. In *Proceedings of FSCD*, volume 167 of *LIPIcs*, pages 34:1–34:21, 2020. URL: `https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.FSCD.2020.34`, `doi:10.4230/LIPIcs.FSCD.2020.34`.

6 Saumya K. Debray and Prateek Mishra. Denotational and operational semantics for prolog. *J. Log. Program.*, 5(1):61–91, March 1988. `doi:10.1016/0743-1066(88)90007-6`.

7 Cvetan Dunchev, Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. ELPI: fast, embeddable, λProlog interpreter. In *Proceedings of LPAR*, volume 9450 of *LNCS*, pages 460–468. Springer, 2015. URL: `https://inria.hal.science/hal-01176856v1`, `doi:10.1007/978-3-662-48899-7\_32`.

8 Davide Fissore and Enrico Tassi. A new Type-Class solver for Coq in Elpi. In *The Coq Workshop*, July 2023. URL: `https://inria.hal.science/hal-04467855`.

9 Davide Fissore and Enrico Tassi. Higher-order unification for free!: Reusing the meta-language unification for the object language. In *Proceedings of PPDP*, pages 1–13. ACM, 2024. `doi:10.1145/3678232.3678233`.

10 Davide Fissore and Enrico Tassi. Determinacy checking for elpi: an higher-order logic programming language with cut. In *Practical Aspects of Declarative Languages: 28th International*

Symposium, PADL 2026, Rennes, France, January 12–13, 2026, Proceedings, pages 77–95, Berlin, Heidelberg, 2026. Springer-Verlag. `doi:10.1007/978-3-032-15981-6_5`.

11 Benjamin Grégoire, Jean-Christophe Léchenet, and Enrico Tassi. Practical and sound equality tests, automatically. In *Proceedings of CPP*, page 167–181. Association for Computing Machinery, 2023. `doi:10.1145/3573105.3575683`.

12 Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. Implementing type theory in higher order constraint logic programming. In *Mathematical Structures in Computer Science*, volume 29, pages 1125–1150. Cambridge University Press, 2019. `doi:10.1017/S0960129518000427`.

13 Robbert Krebbers, Luko van der Maas, and Enrico Tassi. Inductive Predicates via Least Fixpoints in Higher-Order Separation Logic. In Yannick Forster and Chantal Keller, editors, *16th International Conference on Interactive Theorem Proving (ITP 2025)*, volume 352 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 27:1–27:21, Dagstuhl, Germany, 2025. Schloss Dagstuhl – Leibniz-Zentrum für Informatik. URL: `https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.ITP.2025.27`, `doi:10.4230/LIPIcs.ITP.2025.27`.

14 Dale Miller. A logic programming language with lambda-abstraction, function variables, and simple unification. In *Extensions of Logic Programming*, pages 253–281. Springer, 1991.

15 Dale Miller and Gopalan Nadathur. *Programming with Higher-Order Logic*. Cambridge University Press, 2012.

16 Cornelia Pusch. Verification of compiler correctness for the wam. In Gerhard Goos, Juris Hartmanis, Jan van Leeuwen, Joakim von Wright, Jim Grundy, and John Harrison, editors, *Theorem Proving in Higher Order Logics*, pages 347–361, Berlin, Heidelberg, 1996. Springer Berlin Heidelberg.

17 Enrico Tassi. Elpi: an extension language for Coq (Metaprogramming Coq in the Elpi λProlog dialect). In *The Fourth International Workshop on Coq for Programming Languages*, January 2018. URL: `https://inria.hal.science/hal-01637063`.

18 Enrico Tassi. Deriving proved equality tests in Coq-Elpi. In *Proceedings of ITP*, volume 141 of *LIPIcs*, pages 29:1–29:18, September 2019. URL: `https://inria.hal.science/hal-01897468`, `doi:10.4230/LIPIcs.CVIT.2016.23`.

19 Luko van der Maas. Extending the Iris Proof Mode with inductive predicates using Elpi. Master's thesis, Radboud University Nijmegen, 2024. `doi:10.5281/zenodo.12568604`.

20 David H.D. Warren. An Abstract Prolog Instruction Set. Technical Report Technical Note 309, SRI International, Artificial Intelligence Center, Computer Science and Technology Division, Menlo Park, CA, USA, October 1983. URL: `https://www.sri.com/wp-content/uploads/2021/12/641.pdf`.