

Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

Jane Open Access  

Dummy University Computing Laboratory, [optional: Address], Country

My second affiliation, Country

Joan R. Public¹  

Department of Informatics, Dummy College, [optional: Address], Country

Abstract

9 Lorem ipsum dolor sit amet, consectetur adipiscing elit. Praesent convallis orci arcu, eu mollis dolor.
10 Aliquam eleifend suscipit lacinia. Maecenas quam mi, porta ut lacinia sed, convallis ac dui. Lorem
11 ipsum dolor sit amet, consectetur adipiscing elit. Suspendisse potenti.

12 **2012 ACM Subject Classification** Replace ccsdesc macro with valid one

13 **Keywords and phrases** Dummy keyword

14 **Digital Object Identifier** 10.4230/LIPIcs.CVIT.2016.23

15 **Funding Jane Open Access:** (Optional) author-specific funding acknowledgements

16 **Joan R. Public:** [funding]

17 **Acknowledgements** I want to thank ...

1 Introduction

19 ELPI is a dialect of λ PROLOG (see [14, 15, 7, 12]) used as an extension language for the ROCQ
20 prover (formerly the Coq proof assistant). ELPI has become an important infrastructure
21 component: several projects and libraries depend on it [13, 3, 4, 19, 8, 9]. Examples include
22 the Hierarchy-Builder library-structuring tool [5] and Derive [17, 18, 11], a program-and-proof
23 synthesis framework with industrial applications at SkyLabs AI.

24 Starting with version 3, ELPI gained a static analysis for determinacy [10] to help users
25 tame backtracking. ROCQ users are familiar with functional programming but not necessarily
26 with logic programming and uncontrolled backtracking is a common source of inefficiency
27 and makes debugging harder. The determinacy checkers identifies predicates that behave
28 like functions, i.e., predicates that commit to their first solution and leave no *choice points*
29 (places where backtracking could resume).

30 This paper reports our first steps towards a mechanization, in the ROCQ prover, of the
31 determinacy analysis from [10]. We focus on the control operator *cut*, which is useful to
32 restrict backtracking but makes the semantic depart from a pure logical reading.

33 We formalize two operational semantics for PROLOG with cut. The first is a stack-
34 based semantics that closely models ELPI's implementation and is similar to the semantics
35 mechanized by Pusch in ISABELLE/HOL [16] and to the model of Debray and Mishra [6,
36 Sec. 4.3]. This stack-based semantics is a good starting point to study further optimizations
37 used by standard PROLOG abstract machines [20, 1], but it makes reasoning about the scope
38 of *cut* difficult. To address that limitation we introduce a tree-based semantics in which the
39 branches pruned by *cut* are explicit and we prove the two semantics equivalent. Using the

¹ Optional footnote, e.g. to mark corresponding author

23:2 Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

```

Inductive P := IP of nat. Inductive D := ID of nat. Inductive V := IV of nat.

Inductive Tm :=
| Tm_P of P      | Tm_D   of D      | Tm_V of V      | Tm_App of Tm & Tm.

Inductive Callable :=
| Callable_P of P | Callable_App of Callable & Tm.

```

Figure 1 Tm and Callable types

40 tree-based semantics we then show that if every rule of a predicate passes the determinacy
41 analysis, the call to a deterministic predicate does not leave any choice points.

2 Common code: the language

put unif and progs
gram in variables
hides from types
46 Before going to the two semanticcs, we show the piece of data structure that are shared by
the them. The smallest unit of code that we can use in the langauge is an atom. The atom
inductive (see Type 1) is either a cut or a call. A call carries a callable term (see Figure 1).
A term (Tm) is either a predicate, a datum, a variable or the binary application of a term to
another. A Callable is a term accepting predicates only predicates as functors.

```

48 Inductive A := cut | call : Callable -> A.                               (1)
49 Record R := mkR { head : Callable; premises : list A }.                  (2)
50 Record program := { rules : seq R; sig : sigT }.                         (3)
51 Definition Sigma := {fmap V -> Tm}.                                       (4)
52 Definition bc : Unif -> program -> fvS -> Callable ->
      Sigma -> (fvS * seq (Sigma * R)) :=                                (5)

```

53 A rule (see Type 2) is made a head of type term and a list of premises, the premises are
54 atoms. A program (see Type 3) is made by a list of rules and a mapping from predicates to
55 their signatures. The type sigT is the classic type from the simply typed lambda calculus, i.e.
56 it is either a base type or an arrow. We decorate arrows to know the mode of the lhs type.

57 A substitution (see Type 4) is a mapping from variables to terms. It is the output of a
58 successful query and is often called the output of a query.

```

Record Unif := {
  unify : Tm -> Tm -> Sigma -> option Sigma;
  matching : Tm -> Tm -> Sigma -> option Sigma;
}.

```

59 The backchain function (bc, see Type 5) filters the rules in the program that can be
60 used on a given query. It takes: a unificator U which explains how to unify terms up to
61 standard unification (for output terms) or matching (for input terms); a program P to explore
62 and filter; a set S of free variable (fvS) allowing to fresh the program P by renaming the
63 its variables; a query q ; and the substitution σ in which the query q lives. The result of a
64 backchain operation is couple made of an extension of S containing the new variales that
65 have been allocated during the unification phase and a list of filtered rules r accompagnate
66 by their a substition. This substitution is the result of the unification of q with the head of
67 each rule in r .

68 In Figure 2, we have an example of a simple ELPI program which will be used in the
69 following section of the paper as an example to show how backtracking and the cut operator
70 works in the semantcis we propose. The translation of these rules in the ROCQ representation
71 is straightforward.

```
f 1 2.    f 2 3.    r 2 4.    r 2 8.
g X X.          % r1
g X Z :- r X Z, !.   % r2
g X Z :- f X Y, f Y Z.   % r3
```

Figure 2 Small ELPI program example

72 2.1 The cut operator

73 The semantics of the cut operator adopted in the ELPI language corresponds to the *hard*
 74 *cut* operator of standard SWI-PROLOG. This operator has two primary purposes. First,
 75 it eliminates all alternatives that are created either simultaneously with, or after, the
 76 introduction of the cut into the execution state.

77 To illustrate this high-level description, consider the program shown in Figure 2 and the
 78 query $q = g 2 Z$. All three rules for g can be used on the query q . They are tried according
 79 to their order of appearance in the program: rule r_1 is tried first, followed by r_2 , and r_3 .

80 The first rule has no premises and immediately returns the assignment $Z = 2$. However,
 81 the computation does not terminate at this point, since two additional unexplored alternatives
 82 remain, corresponding to the premises of rules r_2 and r_3 .

83 The premises of rule r_2 are $r 2 Z, !$. At this stage, the role of the cut becomes apparent.
 84 If the premise $r 2 Z$ succeeds, the cut commits to this choice and removes the premises of rule
 85 r_3 from the alternative list, as they were generated at the same point as the cut. Moreover,
 86 if the call $r 2 Z$ itself produces multiple alternatives, only the first one is committed, while
 87 the remaining alternatives are discarded. This is because such alternatives have been created
 88 at a deeper depth in the search tree than the cut.

89 Concretely, the call $r 2 Z$ yields two solutions, assigning Z the values 4 and 8, respectively.
 90 The second solution is eliminated by the cut, and only the first assignment is preserved.

91 3 Semantics intro

92 We propose two operational semantics for a logic program with cut. The two semantics are
 93 based on different syntaxes, the first syntax (called tree) exploits a tree-like structure and is
 94 ideal both to have a graphical view of its evolution while the state is being interpreted and
 95 to prove lemmas over it. The second syntax, called elpi, is the ELPI's syntax and has the
 96 advantage of reducing the computational cost of cutting and backtracking alternatives by
 97 using shared pointers. We aim to prove the equivalence of the two semantics together with
 98 some interesting lemmas of the cut behavior.

99 3.1 Tree semantics

```
Inductive tree :=
| KO | OK | TA : A -> tree
| Or : option tree -> Sigma -> tree -> tree
| And : tree -> seq A -> tree -> tree.
```

100 In the tree we distinguish 5 main cases: *KO*, *OK*, and are special meta-symbols representing,
 101 respectively, the failed and a successful terminal. These symbols are considered meta
 102 because they are internal intermediate symbols used to give structure to the tree.

103 The *TA* constructor (acronym for tree-atom) is the constructor of atoms in the tree.

23:4 Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis

```

Fixpoint path_end A :=
  match A with
  | OK | KO | TA _ => A
  | Or None _ B => path_end B
  | Or (Some A) _ _ => path_end A
  | And A _ B =>
    match path_end A with
    | OK => path_end B
    | A => A
  end
end.

```

(a) Definition of *path_end*

104 The two recursive cases of a tree are the *Or* and *And* non-terminals. The *Or* non-terminal
 105 $A \vee B_\sigma$ denotes a disjunction between two trees A and B . The first branch is optional, if
 106 absent it represents a dead tree, i.e. a tree that has been entirely. The second branch is
 107 annotated with a suspended substitution σ so that, upon backtracking to B , σ is used as the
 108 initial substitution for the execution of B .

109 The *And* non-terminal $A \wedge_{B_0} B$ represents a conjunction of two trees A and B . We call B_0
 110 the reset point for B ; it is used to restore the state of B to its initial form if a backtracking
 111 operation occurs on A . Intuitively, let $t2l$ be the function flattening a tree in a list of sequents
 112 disjunction, in PROLOG-like syntax the tree $A \wedge_{B_0} B$ becomes $(A_1, t2l B); (A_2, B_0); \dots; (A_n, B_0)$
 113 where $t2l(A) = A_1, \dots, A_n$.

114 A graphical representation of a tree is shown in Figure 4a. To make the graph more
 115 compact, the *And* and *Or* non-terminals are n-ary rather than binary, with right-binding
 116 priority. The *KO* and *Dead* terminals act as the neutral elements in the *Or* list, while *OK* is
 117 the neutral element of the *And* list.

118 The interpretation of a tree is performed by two main routines: *step* and *next_alt* that
 119 traverse the tree depth-first, left-to-right. Then, then *run* inductive makes the transitive
 120 closure of step *step* and *next_alt*: it iterates the calls to its auxiliary functions. In Types 7–9
 121 we give the types contracts of these symbols where *fv* is a set of variable names.

```

122   Inductive step_tag := Expanded | CutBrothers | Failed | Success.           (6)
123   Definition step : program -> fvS -> Sigma -> tree -> (fvS * step_tag * tree) := (7)
124   Definition next_alt : bool -> tree -> option tree :=                      (8)
125   Inductive run (p : program) : fvS -> Sigma -> tree ->                  (9)
                                option Sigma -> option tree -> bool -> fvS -> Prop :=
```

126 A particular tree we want to identify is a *is_dead* tree (defined in ??). This tree has the
 127 property to never produce a solution: it is either the *Dead* tree or both branches of *Or* are
 128 dead, or the lhs of *And* is dead. In the latter case, we note that B can be non-dead, but this
 129 is not a problem since the interpreter can run B only if A is non-dead.

130 The prolog interpreter explores the state in DFS strategy, it finds the “first-to-be-explored”
 131 (ftbe) atom of the tree and then interprets it. In a non-*is_dead* tree, we get the ftbe node
 132 via *path_end*, shown in Figure 3a. The *path_end* is either the tree itself if the tree is a leaf.
 133 Otherwise, if the tree is a disjunction, the path continues on the left- or the right-subtree
 134 depending of if the the lhs is a *is_dead* tree. In the *Or* case we are clearing ignoring the
 135 dead (ghost) state.

136 In the case of a conjunction, it is more interesting to see what happens. If the *path_end*
 137 p of the lhs is a success then we look for the *path_end* in the rhs, otherwise we return p . In

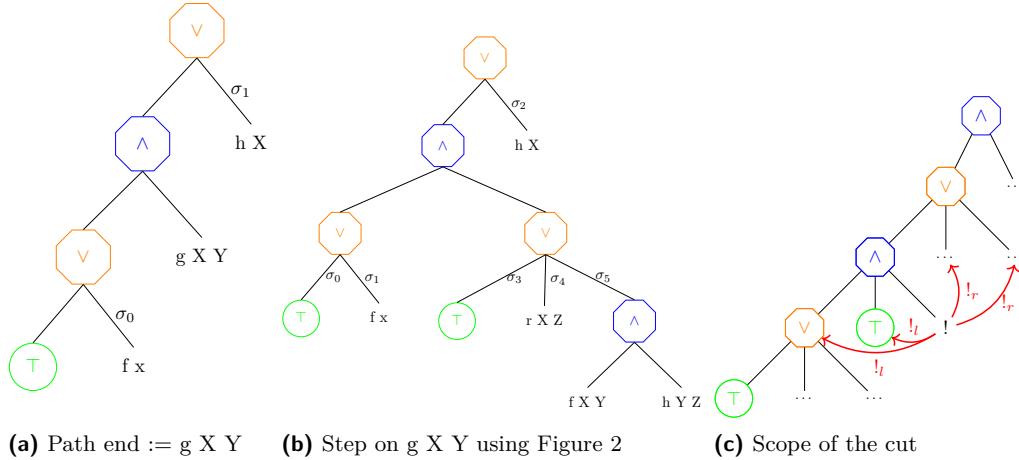


Figure 4 Some tree representations

138 Figure 4a the *path_end* of the tree is *g X*.

139 Below we define two special kind of trees depending on their pathend.

140 **Definition** *success A* := *path_end A == OK*. (1)

141 **Definition** *failed A* := (*path_end A == KO*). (2)

142 3.1.1 The *step* procedure

143 The *step* procedure takes as input a program, a set of free variables (fv), a substitution, and
144 a tree, and returns an updated set of free variables, a *step_tag*, and an updated tree.

145 Free variables are those variables that appear in a tree; they are used in the backchaining
146 operation to refresh the variables in the program.

147 The *step_tag* indicates the type of internal tree step that has been performed. **CutBrothers**
148 denotes the interpretation of a superficial cut, i.e., a cut whose parent nodes are all *And*-nodes.
149 **Expanded** denotes the interpretation of non-superficial cuts or predicate calls. **Failure** and
150 **Success** are returned for, respectively, *failed* and *success* trees.

151 The step procedure is intended to interpretate atoms, that is, it returns the identity for
152 *success* and *failed* tree.

153 **Lemma** *success_step u p fv s A*: *success A* \rightarrow *step u p fv s A = (fv, Success, A)*. (1)
154 **Lemma** *failed_step u p fv s1 A*: *failed A* \rightarrow *step u p fv s1 A = (fv, Failed, A)*. (2)

155 Therefore, *step* produces interesting results if the path-end of the input tree is either a
156 call or a cut.

157 *Call step* The interpretation of a call *c* stars by calling the *bc* function on *c*. The output
158 list *l* is taken to represent build the new subtree. If *l* is empty then *KO* tree is returned,
159 otherwise the subtree is a right-skewed tree made of *n* inner *Or* nodes, where *n* is the length
160 of *l*. The root has *KO* as left child. The lhs of the other nodes is a right-skewed tree of *And*
161 nodes. The *And* nodes are again a right-seked tree containing premises of the selected rule .

if we go right
in the tree, the
subst is the one
in the or...
dire dei reset
point

162 A step in the tree of Figure 4a makes a backchain operation over the query *g X Y* and, in
163 the program defined in Figure 2, the new tree would be the one in Figure 4b. We have put a
164 red border around the new generated subtree. It is a disjunction of four subtrees: the first
165 node is the *KO* node (by default), the second is *OK*, since *r1* has no premises, the third and
166 the fourth contains the premises of respectively *r2* and *r3*.

dire che le sostituzioni del backchain sono importanti

167 *Cut step* The latter case is delicate since interpreting a cut in a tree has three main
168 impacts: at first it is replaced by the *OK* node, then some special subtrees, in the scope
169 of the *Cut*, are cut away: in particular we need to soft-kill the left-siblings of the *Cut* and
170 hard-kill the right-uncles of the the *Cut*.

171 ▶ **Definition 1** (Left-siblings (resp. right-sibling)). *Given a node A, the left-siblings (resp.*
172 *right-sibling) of A are the list of subtrees sharing the same parent of A and that appear on*
173 *its left (resp. right).*

174 ▶ **Definition 2** (Right-uncles). *Given a node A, the right-uncles of A are the list of right-sibling*
175 *of the father of A.*

176 ▶ **Definition 3** (Soft-kill, $!_l$). *Given a successfull tree t, soft-kill replaces all the leaves of the*
177 *tree with the node KO except for the path in t leading to the OK node.*

178 ▶ **Definition 4** (Hard-kill, $!_r$). *Given a tree t, hard-kill replaces all the leaves of the tree with*
179 *the node KO*

180 An example of the impact of the cut is show in Figure 4c. The step routine interprets
181 the cut since it is the node in its path-end. In the example we have 4 arrow tagged with the
182 $!_l$ or $!_r$ symbols. The $!_l$ arrows go left and soft-kill the pointed subtree, in particular, we can
183 note that both pointed subtree have a success node, this is beacuse, in order to evaluate the
184 cut in the figure, we need a successful path leading to it. The $!_l$ procedure will keep the two
185 *OK* nodes since they are essential to reach the cut, and will kill all the leaves in the other
186 subtrees, for those specific subtrees, $!_l$ behaves as $!_r$. The $!_r$ procedure, instead, immediately
187 starts by removing all leaves in the trees pointed by the red arrows.

dire che step
non aggiunge
mai nuovi dead

188 3.1.2 The *next_alt* procedure

189 It is evident that the *step* alone is not sufficient to reproduce entirely the behavior of the
190 full ELPI solver. In particular, *step* does not perform any backtracking at all: it does not
191 backtrack neither for failures, nor for success, from Lemmas 1 and 2, *step* returns the identity.
192 To do so, we have the *next_alt* procedure: its signature is provided in Type 8 and its
193 implementation in Figure 5.

194 The *next_alt* procedure takes a boolean and a tree and return a new tree if it still contains
195 an alternative. The intuition of *next_alt* is to introduce trasnform failed (or success) path
196 into dead-path by inserting new Dead nodes. The boolean tells if there success leaves should
197 be

198 that is it is allowed to transform *OK* or *KO* leaves into *Dead*, so that the *step* procedure
199 is allowed to ignore the new ghosts states and move on. The boolean taken by *next_alt* tells
200 if it is needed to kill *OK* nodes or not.

201 For example, in Figure 4b the step procedure has created a failed state: its path-end ends
202 in *KO*. The expected behavior of *next_alt* is to take this *KO* node and make it a *Dead*. This
203 allows *step* to continue the exploration of the tree. In particular, the path-end of this new
204 state end in *OK*. The step leaves the state unchanged producing the new substitution. This
205 solution however is not unique, we should be able to backtrack on this successful state. To do
206 so we can call *next_alt* and it will deadify the *OK* node allowing *step* to proceed on r X Z.

207 More concretely the code for *next_alt* is show in

```

Definition next_alt : bool -> tree -> option tree :=
fix next_alt b A :=
  match A with
  | KO => None
  | OK => if b then None else Some OK
  | TA _ => Some A
  | And A BO B =>
    let build_BO A := Some (And A BO (big_and BO)) in
    let reset := obind build_BO (next_alt (success A) A) in
    if success A then
      match next_alt b B with
      | None => reset
      | Some B => Some (And A BO B)
      end
    else if failed A then reset
    else Some (And A BO B)
  | Or A sB B =>
    if A is Some A then
      match next_alt b A with
      | None => obind (fun x => Some (Or None sB x)) (next_alt false B)
      | Some nA => Some (Or (Some nA) sB B)
      end
    else
      omap (fun x => (Or None sB x)) (next_alt b B)
  end.

```

■ **Figure 5** *next_alt* implementation

208 3.1.3 The *run* inductive

209 3.1.4 Valid tree

210 Reasoning on a the tree semantics allows to identify an invariant that

211 3.2 Elpi semantics

212 TODO: dire che la semantica ad albero è più facile per le prove

213 The ELPI interpreter is based on an operational semantics close to the one picked by
214 Pusch in [16], in turn closely related to the one given by Debray and Mishra in [6, Section
215 4.3]. Push mechanized the semantics in Isabelle/HOL together with some optimizations that
216 are present in the Warren Abstract Machine [20, 1].217 In these operational semantics we need to decorate the cut atom with a list of alternative,
218 morally a pointer to a sub-list of the overall alternatives. An atom in the elpi semantics is
219 defined as follows:

```
Inductive alts :=
| no_alt
| more_alt : (Sigma * goals) -> alts -> alts
with goals :=
| no_goals
| more_goals : (A * alts) -> goals -> goals .
```

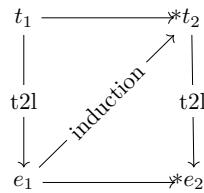
220 We are completely loosing the tree structure. There are no clean reset points. The
221 backtracking operation is simpler: it is the tail function. The cutr and cndl operations
222 disappears: the alternatives are stored directly in the cutE terminal.

223 The elpi interpreter is as follows:

```
(*TODO: add system of rules*)
Inductive nur : Sigma -> goals -> alts -> Sigma -> alts -> Type :=
| StopE s a : nur s nilC a s a
| CutE s s1 a ca r gl : nur s gl ca s1 r -> nur s ((cutE ca) :: gl) a s1 r
| CallE p s s1 a b bs gl r t :
  F u p t s = [:: b & bs] ->
    nur b.1 (save_goals a gl (a2gs1 p b)) (save_alts a gl ((aa2gs p) bs) ++ a) s1 r ->
    nur s ((callE p t) :: gl) a s1 r
| FailE p s s1 s2 t gl a al r :
  F u p t s = [::] -> nur s1 a al s2 r -> nur s ((callE p t) :: gl) ((s1, a) :: al) s2 r.
```

224 The translation of a tree to a list is as follows:

```
Fixpoint t2l (A: tree) s (bt : alts) : alts :=
match A with
| OK           => [:: (s, [::])]
| KO           => [::]
| TA a         => [:: (s, [:: (a,[::])])]
| Or A s1 B   =>
  let 1B := t2l B s1 [::] in
  let 1A := if A is Some A then t2l A s 1B else [::] in
  add_ca_deep bt (1A ++ 1B)
| And A B0 B  =>
  let 1B0 : goals := r2l B0 in
  let 1A  := t2l A s bt in
```



■ **Figure 6** Induction scheme for Theorem 6

```

if lA is [:: (s1A, x) & xs] then
  let xz := add_deepG bt lB0 x in
  let xs := add_deep bt lB0 xs in
  let xs := map (catr lB0) xs in
  let lB := t2l B s1A (xs ++ bt) in
  (map (catl xz) lB) ++ xs
else []
end.
  
```

► **Theorem 5** (tree_to_elpi).

```

225    $\forall A \sigma_1 B \sigma_2 b \sigma_0, \text{vt } A \rightarrow$ 
226    $\text{run}_u \sigma_1 A (\text{Some } \sigma_2) B b \rightarrow$ 
227    $\exists x \text{ xs}, \text{t2l } A \sigma_1 \emptyset = x :: \text{ xs} \wedge \text{nur}_u x.1 \text{ xs } \sigma_2 (\text{t2l } B \sigma_0 \emptyset).$ 
  
```

► **Theorem 6** (elpi_to_tree).

```

228    $\forall \sigma_1 \sigma_2 a na g,$ 
229    $\text{nur}_u \sigma_1 g a \sigma_2 na \rightarrow$ 
230    $\forall \sigma_0 t, \text{vt } t \rightarrow (\text{t2l } t \sigma_0 \emptyset) = ((\sigma_1, g) :: a) \rightarrow$ 
231    $\exists t' n, \text{run}_u \sigma_0 t (\text{Some } \sigma_2) t' n \wedge \text{t2l } t' \sigma_0 \emptyset = na.$ 
  
```

232 The proof of Theorem 6 is based on the idea explained in [2, Section 3.3]. An ideal
 233 statement for this lemma would be: given a function $12t$ transforming an elpi state to a tree,
 234 we would have have that the the execution of an elpi state e is the same as executing run on
 235 the tree resulting from $12t(e)$. However, it is difficult to retrive the strucuture of an elpi state
 236 and create a tree from it. This is because, in an elpi state, we have no clear information
 237 about the scope of an atom inside the list and, therefore, no evident clue about where this
 238 atom should be place in the tree.

239 Our theorem states that, starting from a valid state t which translates to a list of
 240 alternatives $(\sigma_1, g) :: a$. If we run in elpi the list of alternatives, then the execution of the
 241 tree t returns the same result as the execution in elpi. The proof is performed by induction
 242 on the derivations of the elpi execution. We have 4 derivations.

243 We have 4 case to analyse:

244 — **References** —

- 245 1 Hassan Aït-Kaci. *Warren's Abstract Machine: A Tutorial Reconstruction*. The MIT Press, 08
 246 1991. doi:10.7551/mitpress/7160.001.0001.
- 247 2 Yves Bertot. A certified compiler for an imperative language. Technical Report RR-3488,
 248 INRIA, September 1998. URL: <https://inria.hal.science/inria-00073199v1>.

- 249 3 Valentin Blot, Denis Cousineau, Enzo Crance, Louise Dubois de Prisque, Chantal Keller,
 250 Assia Mahboubi, and Pierre Vial. Compositional pre-processing for automated reasoning in
 251 dependent type theory. In Robbert Krebbers, Dmitriy Traytel, Brigitte Pientka, and Steve
 252 Zdancewic, editors, *Proceedings of the 12th ACM SIGPLAN International Conference on
 253 Certified Programs and Proofs, CPP 2023, Boston, MA, USA, January 16-17, 2023*, pages
 254 63–77. ACM, 2023. doi:10.1145/3573105.3575676.
- 255 4 Cyril Cohen, Enzo Crance, and Assia Mahboubi. Trocq: Proof transfer for free, with or
 256 without univalence. In Stephanie Weirich, editor, *Programming Languages and Systems*, pages
 257 239–268, Cham, 2024. Springer Nature Switzerland.
- 258 5 Cyril Cohen, Kazuhiko Sakaguchi, and Enrico Tassi. Hierarchy Builder: Algebraic hierarchies
 259 Made Easy in Coq with Elpi. In *Proceedings of FSCD*, volume 167 of *LIPICS*, pages 34:1–34:21,
 260 2020. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.FSCD.2020.34>.
 261 doi:10.4230/LIPIcs.FSCD.2020.34.
- 262 6 Saumya K. Debray and Prateek Mishra. Denotational and operational semantics for prolog. *J.
 263 Log. Program.*, 5(1):61–91, March 1988. doi:10.1016/0743-1066(88)90007-6.
- 264 7 Cvetan Dunchev, Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. ELPI: fast,
 265 embeddable, λ Prolog interpreter. In *Proceedings of LPAR*, volume 9450 of *LNCS*, pages
 266 460–468. Springer, 2015. URL: <https://inria.hal.science/hal-01176856v1>, doi:10.1007/
 267 978-3-662-48899-7_32.
- 268 8 Davide Fissore and Enrico Tassi. A new Type-Class solver for Coq in Elpi. In *The Coq
 269 Workshop*, July 2023. URL: <https://inria.hal.science/hal-04467855>.
- 270 9 Davide Fissore and Enrico Tassi. Higher-order unification for free!: Reusing the meta-
 271 language unification for the object language. In *Proceedings of PPDP*, pages 1–13. ACM, 2024.
 272 doi:10.1145/3678232.3678233.
- 273 10 Davide Fissore and Enrico Tassi. Determinacy checking for elpi: an higher-order logic program-
 274 ming language with cut. In *Practical Aspects of Declarative Languages: 28th International
 275 Symposium, PADL 2026, Rennes, France, January 12–13, 2026, Proceedings*, pages 77–95,
 276 Berlin, Heidelberg, 2026. Springer-Verlag. doi:10.1007/978-3-032-15981-6_5.
- 277 11 Benjamin Grégoire, Jean-Christophe Léchenet, and Enrico Tassi. Practical and sound equality
 278 tests, automatically. In *Proceedings of CPP*, page 167–181. Association for Computing
 279 Machinery, 2023. doi:10.1145/3573105.3575683.
- 280 12 Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. Implementing type theory
 281 in higher order constraint logic programming. In *Mathematical Structures in Computer
 282 Science*, volume 29, pages 1125–1150. Cambridge University Press, 2019. doi:10.1017/
 283 S0960129518000427.
- 284 13 Robbert Krebbers, Luko van der Maas, and Enrico Tassi. Inductive Predicates via Least
 285 Fixpoints in Higher-Order Separation Logic. In Yannick Forster and Chantal Keller, editors,
 286 *16th International Conference on Interactive Theorem Proving (ITP 2025)*, volume 352 of *Leib-
 287 niz International Proceedings in Informatics (LIPIcs)*, pages 27:1–27:21, Dagstuhl, Germany,
 288 2025. Schloss Dagstuhl – Leibniz-Zentrum für Informatik. URL: [https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.ITP.2025.27](https://drops.dagstuhl.

 289 de/entities/document/10.4230/LIPIcs.ITP.2025.27), doi:10.4230/LIPIcs.ITP.2025.27.
- 290 14 Dale Miller. A logic programming language with lambda-abstraction, function variables, and
 291 simple unification. In *Extensions of Logic Programming*, pages 253–281. Springer, 1991.
- 292 15 Dale Miller and Gopalan Nadathur. *Programming with Higher-Order Logic*. Cambridge
 293 University Press, 2012.
- 294 16 Cornelia Pusch. Verification of compiler correctness for the wam. In Gerhard Goos, Juris
 295 Hartmanis, Jan van Leeuwen, Joakim von Wright, Jim Grundy, and John Harrison, editors,
 296 *Theorem Proving in Higher Order Logics*, pages 347–361, Berlin, Heidelberg, 1996. Springer
 297 Berlin Heidelberg.
- 298 17 Enrico Tassi. Elpi: an extension language for Coq (Metaprogramming Coq in the Elpi λ Prolog
 299 dialect). In *The Fourth International Workshop on Coq for Programming Languages*, January
 300 2018. URL: <https://inria.hal.science/hal-01637063>.

- 301 **18** Enrico Tassi. Deriving proved equality tests in Coq-Elpi. In *Proceedings of ITP*, volume 141 of
302 *LIPICS*, pages 29:1–29:18, September 2019. URL: <https://inria.hal.science/hal-01897468>,
303 doi:10.4230/LIPIcs.CVIT.2016.23.
- 304 **19** Luko van der Maas. Extending the Iris Proof Mode with inductive predicates using Elpi.
305 Master’s thesis, Radboud University Nijmegen, 2024. doi:10.5281/zenodo.12568604.
- 306 **20** David H.D. Warren. An Abstract Prolog Instruction Set. Technical Report Technical Note 309,
307 SRI International, Artificial Intelligence Center, Computer Science and Technology Division,
308 Menlo Park, CA, USA, October 1983. URL: <https://www.sri.com/wp-content/uploads/2021/12/641.pdf>.