

# **Operational semantics for Prolog with Cut in Rocq and its application to determinacy analysis**

**Jane Open Access**  

Dummy University Computing Laboratory, [optional: Address], Country

My second affiliation, Country

**Joan R. Public<sup>1</sup>**  

Department of Informatics, Dummy College, [optional: Address], Country

## **Abstract**

9 Lorem ipsum dolor sit amet, consectetur adipiscing elit. Praesent convallis orci arcu, eu mollis dolor. Aliquam  
10 eleifend suscipit lacinia. Maecenas quam mi, porta ut lacinia sed, convallis ac dui. Lorem ipsum dolor sit amet,  
11 consectetur adipiscing elit. Suspendisse potenti.

12 **2012 ACM Subject Classification** Replace ccsdesc macro with valid one

13 **Keywords and phrases** Dummy keyword

14 **Digital Object Identifier** 10.4230/LIPIcs.CVIT.2016.23

15 **Funding** Jane Open Access: (Optional) author-specific funding acknowledgements

16 Joan R. Public: [funding]

17 **Acknowledgements** I want to thank ...

## **1 Introduction**

19 Elpi is a dialect of  $\lambda$ Prolog (see [14, 15, 7, 12]) used as an extension language for the Rocq Prover  
20 (formerly the Coq proof assistant) that has become an important piece of infrastructure. Several  
21 projects and libraries depend on Elpi [13, 3, 4, 19, 8, 9], for example the Hierarchy-Builder  
22 library-structuring tool [5], and Derive [17, 18, 11], a program-and-proof synthesis framework with  
23 industrial applications at SkyLabs AI.

24 In version 3 Elpi was equipped with a static analysis for determinacy [10] to tame backtracking.  
25 Rocq users are familiar with functional programming but not necessarily with logic programming  
26 and uncontrolled backtracking is a recurrent source of inefficient and hard-to-debug code. The  
27 static analysis identifies “functions”, i.e., predicates that commit to the first result they generate by  
28 leaving no choice points (opportunities for backtracking).

29 This paper is a first step toward the mechanization in the Rocq Prover of the static analysis  
30 from [10] and it focusses on the control operator `cut`. This operator is both the ally to control  
31 backtracking and the enemy when it comes to describing the semantics of the language that  
32 becomes operational departing from the realm of logic.

33 This paper describes the mechanization of two operational semantics for Prolog. One operational  
34 semantic is based on a stack of choice points and reflects closely the implementation of Elpi. This  
35 semantics is close to the one mechanized by Pusch in Isabelle/HOL [16], in turn closely related to  
36 the one given by Debray and Mishra in [6, Section 4.3]. This semantics is well suited to describe  
37 some optimizations that are present in the standard Prolog abstract machine [20, 1], but is not  
38 amenable to reason about the scope of cut, that is paramount in the study of determinacy. Hence  
39 we introduce a tree-based semantics where the branches cut by the cut operator are explicit and we

<sup>1</sup> Optional footnote, e.g. to mark corresponding author



© Jane Open Access and Joan R. Public;  
licensed under Creative Commons License CC-BY 4.0

42nd Conference on Very Important Topics (CVIT 2016)

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:7



Leibniz International Proceedings in Informatics

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

40 prove it is equivalent to the stack-based one. Finally, using the tree-based semantics we establish  
 41 that predicates where each rule passes the static analysis for determinacy do not leave choice points.

## 42 2 Common code: the language

**Inductive** Tm :=

```
| Tm_Kp      : Kp -> Tm
| Tm_Kd      : Kd -> Tm
| Tm_V       : V   -> Tm
| Tm_Comb    : Tm -> Tm -> Tm.
```

**Inductive** Callable :=

```
| Callable_Kp   : Kp -> Callable
| Callable_V    : V   -> Callable
| Callable_Comb : Callable -> Tm -> Callable.
```

**Inductive** RCallable :=

```
| RCallable_Kp   : Kp -> RCallable
| RCallable_Comb : RCallable -> Tm -> RCallable.
```

43 A callable term is a term without a data constructor as functor.

44 An rcallable is a term with rigid head.

**Inductive** A := **cut** | call : Callable -> A.

45 An atom is the smallest syntactic unit that can be executed in a prolog program  $\mathcal{P}$ .

**Record** R := mkR { head : RCallable; premises : list A }.

46 We exploit the typing system to ensure that the head of a "valid" rule is a term with rigid head.

(\*simpler than in the code: signatures of preds are hidden\*)

**Definition** program := seq R.

47 A program is made by a list of rules. Rules in  $\mathcal{P}$  are indexed by their position in the list. Given a  
 48 list of rules  $\mathcal{R}$  and two indexes  $i$  and  $j$ , s.t.  $i \neq j$  then,  $\mathcal{R}_i$  has a higher priority than  $\mathcal{R}_j$ .

49 Sigma is a substitution mapping variables to their term instantiation.

**Definition** Sigma := {fmap V -> Tm}.

50 The backchaining algorithm is the function  $\mathcal{B}$  aims to filter only the rules in the program  $\mathcal{P}$   
 51 having rules unifying with the current query  $q$  in a given substitution  $\sigma$  using the list of modes  $m$ .  
 52 In particular  $\mathcal{B}$  returns for each selected rule  $r$  a substitution  $\sigma'$  that is the substitution obtained by  
 53 the unification of the query and the head of  $r$ .

$$\mathcal{B} : (\mathcal{P}, \sigma, q) \rightarrow \text{seq}(\sigma * R)$$

## 54 3 Semantics intro

55 We propose two operational semantics for a logic program with cut. The two semantics are based  
 56 on different syntaxes, the first syntax (called tree) exploits a tree-like structure and is ideal to have  
 57 a graphical view of its evaloution while the prorgma is being intepreted. The second syntax is  
 58 the elpi's syntax, we call it therefore elpi. We aim to prove the equivalence of the two semantics  
 59 together with some interesting lemmas of the cut behavior.

### 60 3.1 Tree semantics

```
Inductive tree :=  
| Bot | OK | Dead  
| TA : A -> tree  
| Or : tree -> Sigma -> tree -> tree  
| And : tree -> seq A -> tree -> tree.
```

61 In the tree we distinguish 6 main cases: Bot and OK are respectively the standard fail  $\perp$  and true  
 62  $\top$  predicates of prolog. Dead is a special symbol representing a ghost state, that is, a state useful  
 63 to keep the structure of a tree from an execution to another but that is completely ignored by the  
 64 interpretation of the program.

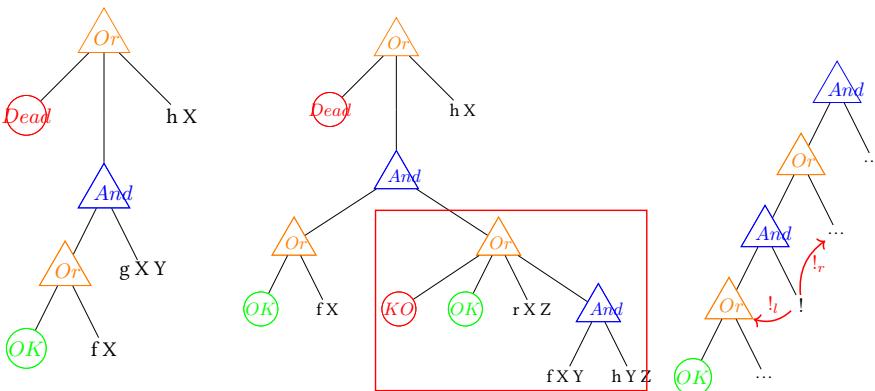
65 TA, standing for tree-atom, is a terminal of the tree containing an atom.

66 The two recursive cases of a tree are the Or and the And non-terminals. The Or non-terminals  
 67  $A \vee B_\sigma$  stands for a disjunction between two trees  $A$  and  $B$ . The second tree branch is decorated  
 68 with a suspended substitution  $\sigma$  so that, when we backtrack to  $B$ , we use  $\sigma$  as initial substitution for  
 69  $B$ .

70 The And non-terminal  $A \wedge_{B_0} B$  represents of a conjunction of two trees  $A$  and  $B$ . We call  
 71  $B_0$  the reset-point for  $B$  and is used to resume the  $B$  state in its initial form if some backtracking  
 72 operation is performed on  $A$ . A graphical tree representation is shown in Figure 1a. For the sake  
 73 of making our graph more compact, the And and Or non-terminals are n-ary (rather than binary),  
 74 with right-binding priority. We are representing the

75 The interpretation of a tree is performed by two main routines: step and next\_alt that  
 76 traverse the tree depth-first, left-to-right.

77 We get the first to-be-explored terminal in the tree by getting the end of a path. This path is  
 78 created from a tree traversal starting from the roots and immediately ends if the tree is not neither a  
 79 disjunction, nor a conjunction: the to-be-explored terminal is the tree itself. Otherwise, if the tree is  
 80 a disjunction, the path continues on the left- or the right-subtree depending of if the path of the lhs  
 81 is a dead node. In the case of a conjunction, we look for the path of the rhs. If this path returns a  
 82 success, we build a path in the rhs, otherwise, we return the lhs. In Figure 1a the first non-explored  
 83 node is g X.



**Figure 1** Tree with first non explored node g X

84 The step procedure takes a tree and explores it using the path strategy. A success (i.e. a tree  
 85 with path ending with OK) and failed tree (i.e. a tree with path ending with KO or Dead) is returned  
 86 as it. The two interesting cases are when the path ends with a call or a cut.

*Call step* In the former case the call node is replaced with a new subtree made by the rules returned by the  $\mathcal{B}$  function. If  $\mathcal{B}$  returns a list  $l$ , if  $l$  is empty then KO tree is returned, otherwise the call is replaced by right-skewed tree made of  $n$  inner Or nodes, where  $n$  is the length of  $l$ . The root Or-node has KO as left child. The lhs of the other nodes is a right-skewed tree of And nodes. The And nodes are again a right-skewed tree containing then atoms (either cut or call) taken from the list  $l$ .

```

g X X. % r1
g X Z :- r X Z. % r2
g X Z :- f X Y, h Y Z. % r3

```

93 A step in the tree in Figure 1a make a backchain operation over the query  $g X Y$  and, in the  
 94 program above, the new tree would be the one in Figure 1b. We have put a red border around  
 95 the new generated subtree. It is a disjunction of four subtrees: the first node is the Dead node (by  
 96 default), the second is OK, since  $r_1$  has no premises, the third and the fourth contains the premises  
 97 of respectively  $r_2$  and  $r_3$ .

*Cut step* The latter case is delicate since interpreting a cut in a tree has three main impacts: at first the cut node is replaced by a OK node, but then we need to cut-away the subtrees that are in the scope of the cut: in particular we need to soft-kill the left-siblings of the Cut and hard-kill the right-uncles of the the Cut.

► **Definition 1** (Left-siblings (resp. right-sibling)). Given a node  $A$ , the left-siblings (resp. right-sibling) of  $A$  are the list of subtrees sharing the same parent of  $A$  and that appear on its left (resp. right).

► **Definition 2** (Right-uncles). *Given a node A, the right-uncles of A are the list of right-sibling of the father of A.*

► **Definition 3 (Soft-kill).** Given a tree  $t$ , soft-kill replaces all the leaves of the tree with the node KO except for the leaves that are part of the path  $p$  of  $t$ .

▶ **Definition 4** (Hard-kill). Given a tree  $t$ , hard-kill replaces all the leaves of the tree with the node  $KO$

An example of the impact of the cut is show in Figure 1c. The step routine interprets the cut if it is at the end of the current path. In the example we have tagged in red the arrow  $!_l$  indicating which sub-trees is soft-killed and  $!_r$  indicated which is sub-trees are to be hard-killed.

The soft-kill operation replace with the KO node all the

### <sup>115</sup> 3.1.1 Valid tree

116 3.2 Elpi semantics

<sup>117</sup> The Elpi interpreter is based on an operational semantics close to the one picked by Pusch in [16],  
<sup>118</sup> in turn closely related to the one given by Debray and Mishra in [6, Section 4.3]. Push mechanized  
<sup>119</sup> the semantics in Isabelle/HOL together with some optimizations that are present in the Warren  
<sup>120</sup> Abstract Machine [20, 1].

In these operational semantics we need to decorate the cut atom with a list of alternative, morally a pointer to a sub-list of the overall alternatives. An atom in the elpi semantics is defined as follows:

**Inductive** alts :=  
| no\_alt

```

| more_alt : (Sigma * goals) -> alts -> alts
with goals := 
| no_goals
| more_goals : (A * alts) -> goals -> goals .

```

123     We are completely loosing the tree structure. There are no clean reset points. The backtracking  
 124 operation is simpler: it is the tail function. The cutr and cutl operations disappears: the alternatives  
 125 are stored directly in the cutE terminal.

126     The elpi interpreter is as follows:

```

(*TODO: add system of rules*)
Inductive nur : Sigma -> goals -> alts -> Sigma -> alts -> Type :=
| StopE s a : nur s nilc a s a
| CutE s s1 a ca r gl : nur s gl ca s1 r -> nur s ((cutE ca) :: gl) a s1 r
| CallE p s s1 a b bs gl r t :
  F u p t s = [:: b & bs] ->
    nur b.1 (save_goals a gl (a2gs1 p b)) (save_alts a gl ((aa2gs p) bs) ++ a
      nur s ((callE p t) :: gl) a s1 r
| FailE p s s1 s2 t gl a al r :
  F u p t s = [::] -> nur s1 a al s2 r -> nur s ((callE p t) :: gl) ((s1, a)

```

127     The translation of a tree to a list is as follows:

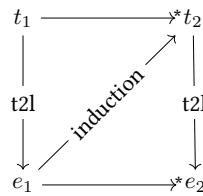
```

Fixpoint t21 (A: tree) s (bt : alts) : alts :=
match A with
| OK => [:: (s, [::])]
| (Bot | Dead) => [::]
| TA a => [:: (s, [:: (a, [::])])]
| Or A s1 B =>
  let 1B := t21 B s1 [::] in
  let 1A := t21 A s 1B in
  add_ca_deep bt (1A ++ 1B)
| And A B0 B =>
  let 1B0 : goals := r21 B0 in
  let 1A := t21 A s bt in
  if 1A is [:: (s1A, x) & xs] then
    let xz := add_deepG bt 1B0 x in
    let xs := add_deep bt 1B0 xs in
    let xs := make_1B0 xs 1B0 in
    let 1B := t21 B s1A (xs ++ bt) in
    (make_1B01 1B xz) ++ xs
  else [::]
end.

```

► **Theorem 5** (tree\_to\_elpi).

128                          $\forall A \sigma_1 B \sigma_2 b \sigma_0, vt A \rightarrow$   
 129                          $run_u \sigma_1 A (Some \sigma_2) B b \rightarrow$   
 130                          $\exists x xs, t21 A \sigma_1 \emptyset = x :: xs \wedge nur_u x.1 x.2 xs \sigma_2 (t21 B \sigma_0 \emptyset).$



■ **Figure 2** Induction scheme for Theorem 6

► **Theorem 6** (elpi\_to\_tree).

131  $\forall \sigma_1 \sigma_2 a na g,$   
 132  $nur_u \sigma_1 g a \sigma_2 na \rightarrow$   
 133  $\forall \sigma_0 t, vt t \rightarrow (t \text{ } 21 t \sigma_0 \emptyset) = ((\sigma_1, g) :: a) \rightarrow$   
 134  $\exists t' n, run_u \sigma_0 t (\text{Some } \sigma_2) t' n \wedge t \text{ } 21 t' \sigma_0 \emptyset = na.$

The proof of Theorem 6 is based on the idea explained in [2, Section 3.3]. An ideal statement for this lemma would be: given a function  $12t$  transforming an elpi state to a tree, we would have have that the execution of an elpi state  $e$  is the same as executing  $\text{run}$  on the tree resulting from  $12t(e)$ . However, it is difficult to retrieve the strucuture of an elpi state and create a tree from it. This is because, in an elpi state, we have no clear information about the scope of an atom inside the list and, therefore, no evident clue about where this atom should be place in the tree.

Our theorem states that, starting from a valid state  $t$  which translates to a list of alternatives  $(\sigma_1, g) :: a$ . If we run in elpi the list of alternatives, then the execution of the tree  $t$  returns the same result as the execution in elpi. The proof is performed by induction on the derivations of the elpi execution. We have 4 derivations.

145 We have 4 case to analyse:

146 — References

- 1 Hassan Ait-Kaci. Warren’s Abstract Machine: A Tutorial Reconstruction. The MIT Press, 08 1991.  
2 Yves Bertot. A certified compiler for an imperative language. Technical Report RR-3488, INRIA, September 1998. URL: <https://inria.hal.science/inria-00073199v1>.

3 Valentin Blot, Denis Cousineau, Enzo Crance, Louise Dubois de Prisque, Chantal Keller, Assia Mahboubi, and Pierre Vial. Compositional pre-processing for automated reasoning in dependent type theory. In Robbert Krebbers, Dmitriy Traytel, Brigitte Pientka, and Steve Zdancewic, editors, Proceedings of the 12th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2023, Boston, MA, USA, January 16-17, 2023, pages 63–77. ACM, 2023. doi:10.1145/3573105.3575676.

4 Cyril Cohen, Enzo Crance, and Assia Mahboubi. Trocq: Proof transfer for free, with or without univalence. In Stephanie Weirich, editor, Programming Languages and Systems, pages 239–268, Cham, 2024. Springer Nature Switzerland.

5 Cyril Cohen, Kazuhiko Sakaguchi, and Enrico Tassi. Hierarchy Builder: Algebraic hierarchies Made Easy in Coq with Elpi. In Proceedings of FSCD, volume 167 of LIPICS, pages 34:1–34:21, 2020. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.FSCD.2020.34>. doi:10.4230/LIPIcs.FSCD.2020.34.

6 Saumya K. Debray and Prateek Mishra. Denotational and operational semantics for prolog. J. Log. Program., 5(1):61–91, March 1988. doi:10.1016/0743-1066(88)90007-6.

- 165 7 Cvetan Dunchev, Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. ELPI: fast, em-  
166 beddable,  $\lambda$ Prolog interpreter. In Proceedings of LPAR, volume 9450 of LNCS, pages 460–  
167 468. Springer, 2015. URL: <https://inria.hal.science/hal-01176856v1>, doi:10.  
168 1007/978-3-662-48899-7\\_32.
- 169 8 Davide Fissoire and Enrico Tassi. A new Type-Class solver for Coq in Elpi. In The Coq Workshop, July  
170 2023. URL: <https://inria.hal.science/hal-04467855>.
- 171 9 Davide Fissoire and Enrico Tassi. Higher-order unification for free!: Reusing the meta-language uni-  
172 fication for the object language. In Proceedings of PPDP, pages 1–13. ACM, 2024. doi:10.1145/  
173 3678232.3678233.
- 174 10 Davide Fissoire and Enrico Tassi. Determinacy checking for elpi: an higher-order logic programming  
175 language with cut. In Practical Aspects of Declarative Languages: 28th International Symposium, PADL  
176 2026, Rennes, France, January 12–13, 2026, Proceedings, pages 77–95, Berlin, Heidelberg, 2026. Springer-  
177 Verlag. doi:10.1007/978-3-032-15981-6\_5.
- 178 11 Benjamin Grégoire, Jean-Christophe Léchenet, and Enrico Tassi. Practical and sound equality tests,  
179 automatically. In Proceedings of CPP, page 167–181. Association for Computing Machinery, 2023.  
180 doi:10.1145/3573105.3575683.
- 181 12 Ferruccio Guidi, Claudio Sacerdoti Coen, and Enrico Tassi. Implementing type theory in higher order  
182 constraint logic programming. In Mathematical Structures in Computer Science, volume 29, pages  
183 1125–1150. Cambridge University Press, 2019. doi:10.1017/S0960129518000427.
- 184 13 Robbert Krebbers, Luko van der Maas, and Enrico Tassi. Inductive Predicates via Least Fixpoints  
185 in Higher-Order Separation Logic. In Yannick Forster and Chantal Keller, editors, 16th International  
186 Conference on Interactive Theorem Proving (ITP 2025), volume 352 of Leibniz International Proceedings  
187 in Informatics (LIPIcs), pages 27:1–27:21, Dagstuhl, Germany, 2025. Schloss Dagstuhl – Leibniz-Zentrum  
188 für Informatik. URL: <https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.ITP.2025.27>, doi:10.4230/LIPIcs.ITP.2025.27.
- 189 14 Dale Miller. A logic programming language with lambda-abstraction, function variables, and simple  
190 unification. In Extensions of Logic Programming, pages 253–281. Springer, 1991.
- 191 15 Dale Miller and Gopalan Nadathur. Programming with Higher-Order Logic. Cambridge University  
192 Press, 2012.
- 193 16 Cornelia Pusch. Verification of compiler correctness for the wam. In Gerhard Goos, Juris Hartmanis, Jan  
194 van Leeuwen, Joakim von Wright, Jim Grundy, and John Harrison, editors, Theorem Proving in Higher  
195 Order Logics, pages 347–361, Berlin, Heidelberg, 1996. Springer Berlin Heidelberg.
- 196 17 Enrico Tassi. Elpi: an extension language for Coq (Metaprogramming Coq in the Elpi  $\lambda$ Prolog dialect). In  
197 The Fourth International Workshop on Coq for Programming Languages, January 2018. URL: <https://inria.hal.science/hal-01637063>.
- 198 18 Enrico Tassi. Deriving proved equality tests in Coq-Elpi. In Proceedings of ITP, volume 141 of LIPIcs,  
199 pages 29:1–29:18, September 2019. URL: <https://inria.hal.science/hal-01897468>,  
200 doi:10.4230/LIPIcs.CVIT.2016.23.
- 201 19 Luko van der Maas. Extending the Iris Proof Mode with inductive predicates using Elpi. Master’s thesis,  
202 Radboud University Nijmegen, 2024. doi:10.5281/zenodo.12568604.
- 203 20 David H.D. Warren. An Abstract Prolog Instruction Set. Technical Report Technical Note 309, SRI  
204 International, Artificial Intelligence Center, Computer Science and Technology Division, Menlo Park,  
205 CA, USA, October 1983. URL: <https://www.sri.com/wp-content/uploads/2021/12/641.pdf>.